

BSTA 551: Statistical Inference

Lesson 1: Introduction to Statistical Inference; Statistics

Jessica Minnier

2026-01-05

Lesson 1: Introduction to Statistical Inference

Welcome to Statistical Inference!

Course Focus: How do we learn about populations from samples?

Welcome to Statistical Inference!

Course Focus: How do we learn about populations from samples?

- **Point Estimation** (Weeks 1-4): What's our best guess for a parameter?

Welcome to Statistical Inference!

Course Focus: How do we learn about populations from samples?

- **Point Estimation** (Weeks 1-4): What's our best guess for a parameter?
- **Confidence Intervals** (Weeks 5-6): What's a plausible range?

Welcome to Statistical Inference!

Course Focus: How do we learn about populations from samples?

- **Point Estimation** (Weeks 1-4): What's our best guess for a parameter?
- **Confidence Intervals** (Weeks 5-6): What's a plausible range?
- **Hypothesis Testing** (Weeks 7-9): Can we make decisions from data?

Welcome to Statistical Inference!

Course Focus: How do we learn about populations from samples?

- **Point Estimation** (Weeks 1-4): What's our best guess for a parameter?
- **Confidence Intervals** (Weeks 5-6): What's a plausible range?
- **Hypothesis Testing** (Weeks 7-9): Can we make decisions from data?
- **Two-Sample Methods** (Week 10): Comparing groups

Today's Goals



Today's Goals

1. Understand the fundamental problem of statistical inference

Today's Goals

1. Understand the fundamental problem of statistical inference
2. Distinguish between populations and samples, parameters and statistics

Today's Goals

1. Understand the fundamental problem of statistical inference
2. Distinguish between populations and samples, parameters and statistics
3. Use R to simulate sampling distributions

Today's Goals

1. Understand the fundamental problem of statistical inference
2. Distinguish between populations and samples, parameters and statistics
3. Use R to simulate sampling distributions
4. Review key properties of expected value and variance

Motivating Example: Clinical Trial

A pharmaceutical company is testing a new blood pressure medication.

The Question: What is the true average reduction in systolic blood pressure?

Motivating Example: Clinical Trial

A pharmaceutical company is testing a new blood pressure medication.

The Question: What is the true average reduction in systolic blood pressure?

What we have: Data from 25 patients in a trial

```
1 # Actual blood pressure reductions (mmHg) from 25 patients
2 bp_reductions <- c(12, 8, 15, 10, 7, 14, 11, 9, 13, 16,
3                   8, 12, 10, 14, 11, 9, 15, 13, 7, 12,
4                   10, 8, 14, 11, 13)
5
6 # What's our estimate of the true mean reduction?
7 mean(bp_reductions)
```

```
[1] 11.28
```

Motivating Example: Clinical Trial

A pharmaceutical company is testing a new blood pressure medication.

The Question: What is the true average reduction in systolic blood pressure?

What we have: Data from 25 patients in a trial

```
1 # Actual blood pressure reductions (mmHg) from 25 patients
2 bp_reductions <- c(12, 8, 15, 10, 7, 14, 11, 9, 13, 16,
3                   8, 12, 10, 14, 11, 9, 15, 13, 7, 12,
4                   10, 8, 14, 11, 13)
5
6 # What's our estimate of the true mean reduction?
7 mean(bp_reductions)
```

```
[1] 11.28
```

This “estimate” is the observed value of the (random variable) sample mean statistic: $\bar{X} = \sum_{i=1}^n X_i$

Motivating Example: Clinical Trial

A pharmaceutical company is testing a new blood pressure medication.

The Question: What is the true average reduction in systolic blood pressure?

What we have: Data from 25 patients in a trial

```
1 # Actual blood pressure reductions (mmHg) from 25 patients
2 bp_reductions <- c(12, 8, 15, 10, 7, 14, 11, 9, 13, 16,
3                   8, 12, 10, 14, 11, 9, 15, 13, 7, 12,
4                   10, 8, 14, 11, 13)
5
6 # What's our estimate of the true mean reduction?
7 mean(bp_reductions)
```


```
[1] 11.28
```

This “estimate” is the observed value of the (random variable) sample mean statistic: $\bar{X} = \sum_{i=1}^n X_i$

But how **reliable** is this estimate? Would we get the same answer with different patients?

Statistical Inference: The Big Picture

The Big Picture: Population vs Sample



Key insight: We use sample data to make inferences about population parameters.

Key Terminology (Devore 6.1)

! Definitions

- **Population:** The entire collection of individuals or measurements of interest
- **Sample:** A subset of the population that we actually observe
- **Parameter:** A numerical characteristic of the population (e.g., μ , σ , p)
- **Statistic:** A numerical characteristic computed from sample data (e.g., $\bar{X}$, S , $\hat{p}$)

Key Terminology (Devore 6.1)

! Definitions

- **Population:** The entire collection of individuals or measurements of interest
- **Sample:** A subset of the population that we actually observe
- **Parameter:** A numerical characteristic of the population (e.g., μ , σ , p)
- **Statistic:** A numerical characteristic computed from sample data (e.g., $\bar{X}$, S , $\hat{p}$)

The Central Challenge: Parameters are fixed but unknown; statistics are known but random.

Parameters vs. Statistics

Concept	Population (Parameter)	Sample (Statistic)
Mean	μ	$\bar{X} = \frac{1}{n} \sum X_i$
Variance	σ^2	$S^2 = \frac{1}{n-1} \sum (X_i - \bar{X})^2$
Proportion	p	$\hat{p} = X/n$
Maximum	θ (upper bound)	$\max(X_1, \dots, X_n)$

Parameters vs. Statistics

Concept	Population (Parameter)	Sample (Statistic)
Mean	μ	$\bar{X} = \frac{1}{n} \sum X_i$
Variance	σ^2	$S^2 = \frac{1}{n-1} \sum (X_i - \bar{X})^2$
Proportion	p	$\hat{p} = X/n$
Maximum	θ (upper bound)	$\max(X_1, \dots, X_n)$

Key Point: Parameters use Greek letters; statistics use Roman letters or “hats.”

The Sampling Process (Devore 6.2)

Random Sampling Assumptions:

1. Each observation X_i is a random variable
2. The X_i are **independent** of each other
3. Each X_i has the **same distribution** (identically distributed)

The Sampling Process (Devore 6.2)

Random Sampling Assumptions:

1. Each observation X_i is a random variable
2. The X_i are **independent** of each other
3. Each X_i has the **same distribution** (identically distributed)

Together: $X_1, X_2, \dots, X_n$ are **i.i.d.** (independent and identically distributed)

The Sampling Process (Devore 6.2)

Random Sampling Assumptions:

1. Each observation X_i is a random variable
2. The X_i are **independent** of each other
3. Each X_i has the **same distribution** (identically distributed)

Together: $X_1, X_2, \dots, X_n$ are **i.i.d.** (independent and identically distributed)

Why This Matters

The i.i.d. assumption allows us to derive the properties of statistics mathematically.

Review: Expected Value and Variance

Review: Expected Value

The **expected value** $E(X)$ is the long-run average of a random variable.

Key Properties We'll Use Today

1. $E(c) = c$ for any constant c
2. $E(cX) = c \cdot E(X)$
3. $E(X + Y) = E(X) + E(Y)$
4. $E\left(\sum_{i=1}^n X_i\right) = \sum_{i=1}^n E(X_i)$

Worked Example: Expected Value of Sample Mean

Problem: If $X_1, X_2, \dots, X_n$ are *iid* observations from a population with mean μ , what is $E(\bar{X})$?

Worked Example: Expected Value of Sample Mean

Problem: If $X_1, X_2, \dots, X_n$ are *iid* observations from a population with mean μ , what is $E(\bar{X})$?

Solution: Let's work through this step by step.

$$E(\bar{X}) = E\left(\frac{1}{n} \sum_{i=1}^n X_i\right)$$

Worked Example: Expected Value of Sample Mean

Problem: If $X_1, X_2, \dots, X_n$ are *iid* observations from a population with mean μ , what is $E(\bar{X})$?

Solution: Let's work through this step by step.

$$\begin{aligned} E(\bar{X}) &= E\left(\frac{1}{n} \sum_{i=1}^n X_i\right) \\ &= \frac{1}{n} E\left(\sum_{i=1}^n X_i\right) \quad (\text{Property 2: constants come out}) \end{aligned}$$

Worked Example: Expected Value of Sample Mean

Problem: If $X_1, X_2, \dots, X_n$ are *iid* observations from a population with mean μ , what is $E(\bar{X})$?

Solution: Let's work through this step by step.

$$\begin{aligned} E(\bar{X}) &= E\left(\frac{1}{n} \sum_{i=1}^n X_i\right) \\ &= \frac{1}{n} E\left(\sum_{i=1}^n X_i\right) \quad (\text{Property 2: constants come out}) \\ &= \frac{1}{n} \sum_{i=1}^n E(X_i) \quad (\text{Property 4: sum of expectations}) \end{aligned}$$

Worked Example: Expected Value of Sample Mean

Problem: If $X_1, X_2, \dots, X_n$ are *iid* observations from a population with mean μ , what is $E(\bar{X})$?

Solution: Let's work through this step by step.

$$\begin{aligned} E(\bar{X}) &= E\left(\frac{1}{n} \sum_{i=1}^n X_i\right) \\ &= \frac{1}{n} E\left(\sum_{i=1}^n X_i\right) \quad (\text{Property 2: constants come out}) \\ &= \frac{1}{n} \sum_{i=1}^n E(X_i) \quad (\text{Property 4: sum of expectations}) \\ &= \frac{1}{n} \cdot n\mu = \mu \quad (\text{Each } E(X_i) = \mu) \end{aligned}$$

Your Turn: Calculate Expected Value

Exercise: A hospital measures the recovery time (in days) for patients after surgery. Let X_1, X_2, X_3 be recovery times for 3 patients. The population mean recovery time is $\mu = 5$ days.

Questions:

1. What is $E(X_1)$?
2. What is $E(X_1 + X_2 + X_3)$?
3. What is $E(\bar{X})$ where $\bar{X} = \frac{X_1 + X_2 + X_3}{3}$?

Your Turn: Calculate Expected Value

Exercise: A hospital measures the recovery time (in days) for patients after surgery. Let X_1, X_2, X_3 be recovery times for 3 patients. The population mean recovery time is $\mu = 5$ days.

Questions:

1. What is $E(X_1)$?
2. What is $E(X_1 + X_2 + X_3)$?
3. What is $E(\bar{X})$ where $\bar{X} = \frac{X_1 + X_2 + X_3}{3}$?

Answers:

1. $E(X_1) = \mu = 5$ days
2. $E(X_1 + X_2 + X_3) = 5 + 5 + 5 = 15$ days
3. $E(\bar{X}) = \frac{15}{3} = 5$ days

Review: Variance

Variance measures the spread of a distribution: $\text{Var}(X) = E[(X - \mu)^2]$

Key Properties We'll Use Today

1. $\text{Var}(c) = 0$ for any constant
2. $\text{Var}(cX) = c^2 \cdot \text{Var}(X)$
3. If X and Y are **independent**: $\text{Var}(X + Y) = \text{Var}(X) + \text{Var}(Y)$

Worked Example: Variance of Sample Mean

Problem: If $X_1, \dots, X_n$ are independent observations with variance σ^2 , what is $\text{Var}(\bar{X})$?

Worked Example: Variance of Sample Mean

Problem: If $X_1, \dots, X_n$ are independent observations with variance σ^2 , what is $\text{Var}(\bar{X})$?

$$\text{Var}(\bar{X}) = \text{Var} \left(\frac{1}{n} \sum_{i=1}^n X_i \right)$$

Worked Example: Variance of Sample Mean

Problem: If $X_1, \dots, X_n$ are independent observations with variance σ^2 , what is $\text{Var}(\bar{X})$?

$$\begin{aligned}\text{Var}(\bar{X}) &= \text{Var} \left(\frac{1}{n} \sum_{i=1}^n X_i \right) \\ &= \frac{1}{n^2} \text{Var} \left(\sum_{i=1}^n X_i \right) \quad (\text{Property 2})\end{aligned}$$

Worked Example: Variance of Sample Mean

Problem: If $X_1, \dots, X_n$ are independent observations with variance σ^2 , what is $\text{Var}(\bar{X})$?

$$\begin{aligned}\text{Var}(\bar{X}) &= \text{Var} \left(\frac{1}{n} \sum_{i=1}^n X_i \right) \\ &= \frac{1}{n^2} \text{Var} \left(\sum_{i=1}^n X_i \right) \quad (\text{Property 2}) \\ &= \frac{1}{n^2} \sum_{i=1}^n \text{Var}(X_i) \quad (\text{Property 3: independence})\end{aligned}$$

Worked Example: Variance of Sample Mean

Problem: If $X_1, \dots, X_n$ are independent observations with variance σ^2 , what is $\text{Var}(\bar{X})$?

$$\begin{aligned}\text{Var}(\bar{X}) &= \text{Var} \left(\frac{1}{n} \sum_{i=1}^n X_i \right) \\ &= \frac{1}{n^2} \text{Var} \left(\sum_{i=1}^n X_i \right) \quad (\text{Property 2}) \\ &= \frac{1}{n^2} \sum_{i=1}^n \text{Var}(X_i) \quad (\text{Property 3: independence}) \\ &= \frac{1}{n^2} \cdot n\sigma^2 = \frac{\sigma^2}{n}\end{aligned}$$

! Key Result

The variance of $\bar{X}$ **decreases** as sample size n increases!

The Sampling Distribution

What is a Sampling Distribution?

Different samples give different values of a statistic. The **sampling distribution** describes this variability.

Sample	Sample Mean ($\bar{x}$)
1	9.90
2	10.31
3	10.03
4	10.85
5	9.17
6	9.31

What is a Sampling Distribution?

Different samples give different values of a statistic. The **sampling distribution** describes this variability.

Sample	Sample Mean ($\bar{x}$)
1	9.90
2	10.31
3	10.03
4	10.85
5	9.17
6	9.31

Each sample gives a **different estimate**, but they cluster around the true value $\mu = 10$.

Simulation: Building a Sampling Distribution

```
1 # Parameters
2 true_effect <- 10 # True mean BP reduction (mmHg)
3 true_sd <- 4      # Standard deviation
4 n_patients <- 25  # Patients per trial
5 n_trials <- 5000  # Number of simulated trials
6
7 # Simulate many clinical trials
8 sampling_distribution <- tibble(trial = 1:n_trials) |>
9   mutate(
10     sample_mean = map_dbl(trial, \(t) {
11       patients <- rnorm(n_patients, true_effect, true_sd)
12       mean(patients)
13     })
14   )
15
16 # Visualize
17 sampling_distribution |>
18   ggplot(aes(x = sample_mean)) +
19   geom_histogram(bins = 50, fill = "steelblue", alpha = 0.7, color = "white") +
20   geom_vline(xintercept = true_effect, color = "red", linewidth = 1.5) +
21   labs(title = "Sampling Distribution of the Sample Mean",
22         subtitle = str_glue("True  $\mu$  = {true_effect}, n = {n_patients}, {n_trials} simulated"),
23         x = "Sample Mean (estimated BP reduction)", y = "Frequency")
```

Simulation: Building a Sampling Distribution

```
1 # Parameters
2 true_effect <- 10 # True mean BP reduction (mmHg)
3 true_sd <- 4      # Standard deviation
4 n_patients <- 25  # Patients per trial
5 n_trials <- 5000  # Number of simulated trials
6
7 # Simulate many clinical trials
8 sampling_distribution <- tibble(trial = 1:n_trials) |>
9   mutate(
10     sample_mean = map_dbl(trial, \(t) {
11       patients <- rnorm(n_patients, true_effect, true_sd)
12       mean(patients)
13     })
14   )
15
16 # Visualize
17 sampling_distribution |>
18   ggplot(aes(x = sample_mean)) +
19   geom_histogram(bins = 50, fill = "steelblue", alpha = 0.7, color = "white") +
20   geom_vline(xintercept = true_effect, color = "red", linewidth = 1.5) +
21   labs(title = "Sampling Distribution of the Sample Mean",
22         subtitle = str_glue("True  $\mu$  = {true_effect}, n = {n_patients}, {n_trials} simulated"),
23         x = "Sample Mean (estimated BP reduction)", y = "Frequency")
```

Simulation: Building a Sampling Distribution

```
1 # Parameters
2 true_effect <- 10 # True mean BP reduction (mmHg)
3 true_sd <- 4      # Standard deviation
4 n_patients <- 25  # Patients per trial
5 n_trials <- 5000  # Number of simulated trials
6
7 # Simulate many clinical trials
8 sampling_distribution <- tibble(trial = 1:n_trials) |>
9   mutate(
10     sample_mean = map_dbl(trial, \(t) {
11       patients <- rnorm(n_patients, true_effect, true_sd)
12       mean(patients)
13     })
14   )
15
16 # Visualize
17 sampling_distribution |>
18   ggplot(aes(x = sample_mean)) +
19   geom_histogram(bins = 50, fill = "steelblue", alpha = 0.7, color = "white") +
20   geom_vline(xintercept = true_effect, color = "red", linewidth = 1.5) +
21   labs(title = "Sampling Distribution of the Sample Mean",
22         subtitle = str_glue("True  $\mu$  = {true_effect}, n = {n_patients}, {n_trials} simulated"),
23         x = "Sample Mean (estimated BP reduction)", y = "Frequency")
```


Simulation: Building a Sampling Distribution

```
1 # Parameters
2 true_effect <- 10 # True mean BP reduction (mmHg)
3 true_sd <- 4      # Standard deviation
4 n_patients <- 25  # Patients per trial
5 n_trials <- 5000  # Number of simulated trials
6
7 # Simulate many clinical trials
8 sampling_distribution <- tibble(trial = 1:n_trials) |>
9   mutate(
10     sample_mean = map_dbl(trial, \(t) {
11       patients <- rnorm(n_patients, true_effect, true_sd)
12       mean(patients)
13     })
14   )
15
16 # Visualize
17 sampling_distribution |>
18   ggplot(aes(x = sample_mean)) +
19   geom_histogram(bins = 50, fill = "steelblue", alpha = 0.7, color = "white") +
20   geom_vline(xintercept = true_effect, color = "red", linewidth = 1.5) +
21   labs(title = "Sampling Distribution of the Sample Mean",
22         subtitle = str_glue("True  $\mu$  = {true_effect}, n = {n_patients}, {n_trials} simulated"),
23         x = "Sample Mean (estimated BP reduction)", y = "Frequency")
```

Simulation: Building a Sampling Distribution

Sampling Distribution of the Sample Mean

True $\mu = 10$, $n = 25$, 5000 simulated trials



Simulation: Seeing Variance Decrease

```
1 # Population parameters
2 true_mean <- 120 # True mean systolic BP
3 true_sd <- 15    # Population standard deviation
4
5 # Simulate sample means for different sample sizes
6 simulation_data <- tibble(n = c(5, 25, 100)) |>
7   cross_join(tibble(sim = 1:2000)) |>
8   mutate(
9     sample_mean = map2_dbl(n, sim, \(size, s) {
10       mean(rnorm(size, true_mean, true_sd))
11     })
12 )
13
14 # Calculate observed standard deviation for each sample size
15 simulation_data |>
16   group_by(n) |>
17   summarize(
18     observed_sd = sd(sample_mean),
19     theoretical_sd = true_sd / sqrt(first(n))
20 )
```

A tibble: 3 × 3

n	observed_sd	theoretical_sd
---	-------------	----------------



Simulation: Seeing Variance Decrease

```
1 # Population parameters
2 true_mean <- 120 # True mean systolic BP
3 true_sd <- 15    # Population standard deviation
4
5 # Simulate sample means for different sample sizes
6 simulation_data <- tibble(n = c(5, 25, 100)) |>
7   cross_join(tibble(sim = 1:2000)) |>
8   mutate(
9     sample_mean = map2_dbl(n, sim, \(size, s) {
10       mean(rnorm(size, true_mean, true_sd))
11     })
12 )
13
14 # Calculate observed standard deviation for each sample size
15 simulation_data |>
16   group_by(n) |>
17   summarize(
18     observed_sd = sd(sample_mean),
19     theoretical_sd = true_sd / sqrt(first(n))
20 )
```

A tibble: 3 × 3

n	observed_sd	theoretical_sd
---	-------------	----------------



11:11

11:11

11:11

Simulation: Seeing Variance Decrease

```
1 # Population parameters
2 true_mean <- 120 # True mean systolic BP
3 true_sd <- 15    # Population standard deviation
4
5 # Simulate sample means for different sample sizes
6 simulation_data <- tibble(n = c(5, 25, 100)) |>
7   cross_join(tibble(sim = 1:2000)) |>
8   mutate(
9     sample_mean = map2_dbl(n, sim, \(size, s) {
10       mean(rnorm(size, true_mean, true_sd))
11     })
12 )
13
14 # Calculate observed standard deviation for each sample size
15 simulation_data |>
16   group_by(n) |>
17   summarize(
18     observed_sd = sd(sample_mean),
19     theoretical_sd = true_sd / sqrt(first(n))
20 )
```

A tibble: 3 × 3

n	observed_sd	theoretical_sd
---	-------------	----------------



11:11

11:11

11:11

Simulation: Seeing Variance Decrease

```
1 # Population parameters
2 true_mean <- 120 # True mean systolic BP
3 true_sd <- 15    # Population standard deviation
4
5 # Simulate sample means for different sample sizes
6 simulation_data <- tibble(n = c(5, 25, 100)) |>
7   cross_join(tibble(sim = 1:2000)) |>
8   mutate(
9     sample_mean = map2_dbl(n, sim, \(size, s) {
10       mean(rnorm(size, true_mean, true_sd))
11     })
12 )
13
14 # Calculate observed standard deviation for each sample size
15 simulation_data |>
16   group_by(n) |>
17   summarize(
18     observed_sd = sd(sample_mean),
19     theoretical_sd = true_sd / sqrt(first(n))
20 )
```

A tibble: 3 × 3

n	observed_sd	theoretical_sd
---	-------------	----------------



11:11


11:11

11:11

Visualizing the Effect of Sample Size

Sampling Distribution of $\bar{X}$ for Different Sample Sizes

Larger $n \rightarrow$ Less spread $\rightarrow$ More precise estimates



Understanding Optimization

Why Optimization Matters in Statistics

Many statistical methods require finding the “best” value of a parameter.

Examples:

- **Maximum Likelihood:** Find the parameter value that makes the observed data most probable
- **Least Squares:** Find the parameter value that minimizes prediction errors
- **Minimum Variance:** Find the estimator with the smallest spread

Why Optimization Matters in Statistics

Many statistical methods require finding the “best” value of a parameter.

Examples:


- **Maximum Likelihood:** Find the parameter value that makes the observed data most probable
- **Least Squares:** Find the parameter value that minimizes prediction errors
- **Minimum Variance:** Find the estimator with the smallest spread

The Problem: How do we find maximums and minimums numerically?

Optimization: The Graphical Intuition

Finding the Maximum of a Function

Where is this function highest?



The maximum occurs where the function reaches its peak.

Concrete Example: Finding the Best Estimate

Scenario: In a clinical trial, 7 out of 10 patients respond to treatment. What's the best estimate of the true response rate p ?

Concrete Example: Finding the Best Estimate

Scenario: In a clinical trial, 7 out of 10 patients respond to treatment. What's the best estimate of the true response rate p ?

Approach: Find the value of p that makes observing “7 out of 10” most likely.

The probability of observing exactly 7 successes in 10 trials is:

$$P(X = 7) = \binom{10}{7} p^7 (1 - p)^3$$

Concrete Example: Finding the Best Estimate

Scenario: In a clinical trial, 7 out of 10 patients respond to treatment. What's the best estimate of the true response rate p ?

Approach: Find the value of p that makes observing “7 out of 10” most likely.

The probability of observing exactly 7 successes in 10 trials is:

$$P(X = 7) = \binom{10}{7} p^7 (1 - p)^3$$

Question: For what value of p is this probability largest?

Grid Search: A Simple Numerical Approach

Idea: Try many values and see which gives the largest result.

```
1 # Try different values of p
2 grid_search <- tibble(p = seq(0.01, 0.99, by = 0.01)) |>
3   mutate(
4     likelihood = dbinom(7, size = 10, prob = p)
5   )
6
7 # Find the maximum
8 grid_search |>
9   slice_max(likelihood, n = 1)
```

```
# A tibble: 1 × 2
  p likelihood
<dbl>      <dbl>
1  0.7      0.267
```

Grid Search: A Simple Numerical Approach

Idea: Try many values and see which gives the largest result.

```
1 # Try different values of p
2 grid_search <- tibble(p = seq(0.01, 0.99, by = 0.01)) |>
3   mutate(
4     likelihood = dbinom(7, size = 10, prob = p)
5   )
6
7 # Find the maximum
8 grid_search |>
9   slice_max(likelihood, n = 1)
```

```
# A tibble: 1 × 2
  p likelihood
<dbl>     <dbl>
1  0.7      0.267
```

The maximum likelihood estimate is $\hat{p} = 0.70 = \frac{7}{10}$. This makes intuitive sense!

Using R's Optimizer

R has built-in functions to find maximums and minimums more precisely:


```
1 # Define the likelihood function
2 likelihood_function <- function(p) {
3   dbinom(7, size = 10, prob = p)
4 }
5
6 # Use optimize() to find the maximum
7 # Note: optimize finds MINIMUM by default, so we negate for maximum
8 result <- optimize(
9   f = function(p) -likelihood_function(p), # Negative to find max
10  interval = c(0, 1)                       # Search between 0 and 1
11 )
12
13 # The maximum occurs at:
14 cat("Maximum likelihood estimate: p =", result$minimum)
```

Maximum likelihood estimate: p = 0.6999843

How Numerical Optimization Works

Numerical Optimization: Searching for the Maximum

The algorithm tries different values, homing in on the peak



Key idea: The algorithm evaluates the function at different points and iteratively narrows in on the maximum.

Your Turn: Numerical Optimization

Exercise: A diagnostic test correctly identifies a disease in 18 out of 25 patients who have it. Find the maximum likelihood estimate for the test's sensitivity p .

```
1 # Fill in the blanks:
2 likelihood_fn <- function(p) {
3   dbinom(____, size = ____, prob = p) # What goes here?
4 }
5
6 result <- optimize(
7   f = function(p) -likelihood_fn(p),
8   interval = c(0, 1)
9 )
10
11 result$minimum # This is the MLE
```

Your Turn: Numerical Optimization

Exercise: A diagnostic test correctly identifies a disease in 18 out of 25 patients who have it. Find the maximum likelihood estimate for the test's sensitivity p .

```
1 # Fill in the blanks:
2 likelihood_fn <- function(p) {
3   dbinom(____, size = ____, prob = p) # What goes here?
4 }
5
6 result <- optimize(
7   f = function(p) -likelihood_fn(p),
8   interval = c(0, 1)
9 )
10
11 result$minimum # This is the MLE
```

```
1 # Solution:
2 likelihood_fn <- function(p) {
3   dbinom(18, size = 25, prob = p)
4 }
5
6 result <- optimize(f = function(p) -likelihood_fn(p), interval = c(0, 1))
7 cat("MLE of sensitivity:", result$minimum) # Should be 18/25 = 0.72
```

MLE of sensitivity: 0.7200103



When Optimization Gets Harder

Sometimes we need to optimize over multiple parameters or complex functions:

```
1 # Example: Finding mean and SD that best fit data
2 patient_data <- c(120, 135, 128, 142, 131, 125, 138, 129, 133, 127)
3
4 # Negative log-likelihood for normal distribution
5 neg_log_lik <- function(params) {
6   mu <- params[1]
7   sigma <- params[2]
8   if (sigma <= 0) return(Inf) # sigma must be positive
9   -sum(dnorm(patient_data, mean = mu, sd = sigma, log = TRUE))
10 }
11
12 # Use optim() for multiple parameters
13 result <- optim(par = c(130, 10), fn = neg_log_lik)
14 cat("MLE for mean:", round(result$par[1], 2), "\n")
```

MLE for mean: 130.8

```
1 cat("MLE for SD:", round(result$par[2], 2), "\n")
```

MLE for SD: 6.13

```
1 cat("Compare to sample mean:", round(mean(patient_data), 2))
```

Compare to sample mean: 130.8

Summary and Looking Ahead

Lesson 1 Summary

Key Concepts:

1. **Statistical Inference:** Using sample data to learn about population parameters

2. **Parameters vs. Statistics:**

- Parameters (μ, σ, p) : Fixed but unknown population values
- Statistics $(\bar{X}, S, \hat{p})$: Calculated from sample data

3. **Sampling Distribution:** The distribution of a statistic across many samples

- $E(\bar{X}) = \mu$ (centered at population mean)
- $\text{Var}(\bar{X}) = \sigma^2/n$ (precision improves with larger n)

4. **Numerical Optimization:** Finding maximum/minimum values

- Grid search: try many values
- `optimize()`: efficient numerical search

Lesson 1 Practice Problems

1. Calculate $E(\bar{X})$ and $\text{Var}(\bar{X})$ for a sample of size $n = 16$ from a population with $\mu = 50$ and $\sigma = 12$.
2. Use `optimize()` to find the MLE for p when you observe 23 successes in 40 trials.
3. A quality control engineer samples 5 items from a production line. If the population mean weight is 100g with SD = 5g, what is the expected value and variance of the sample mean?
4. Simulate the sampling distribution of the sample median for $n = 30$ observations from a Normal(100, 15) distribution. Compare to the sampling distribution of the sample mean.

Next Lesson Preview

Lesson 2: Point Estimation; Bias, Variance, and MSE

- What is a point estimator?
- Bias: systematic error in estimation
- Standard error: precision of estimators
- Mean Squared Error: combining bias and variance
- The bias-variance tradeoff

References

- Devore, Berk, and Carlton. *Modern Mathematical Statistics with Applications* (Springer). Chapters 6.1, 6.2
- Chihara and Hesterberg. *Mathematical Statistics with Resampling and R* (Wiley). Chapter 6.

Questions?

Thank you!