

Lesson 14: P-Values; Power and Sample Size Determination

BSTA 551: Statistical Inference

Jessica Minnier

2026-02-09

Lesson 14: P-Values; Power and Sample Size Determination

Review: Where We Left Off

Lesson 13: Tests About a Population Mean and Proportion

- One-sample z test (σ known) and t test (σ unknown)
- Score test and exact binomial test for proportions
- Seven-step hypothesis testing procedure
- CI-test duality: a $100(1 - \alpha)\%$ CI contains values not rejected at level α

Review: Where We Left Off

Lesson 13: Tests About a Population Mean and Proportion

- One-sample z test (σ known) and t test (σ unknown)
- Score test and exact binomial test for proportions
- Seven-step hypothesis testing procedure
- CI-test duality: a $100(1 - \alpha)\%$ CI contains values not rejected at level α

Today's Goals:

1. Understand P-values as a continuous measure of evidence against H_0
2. Prove that P-values are Uniform(0,1) under H_0 using the probability integral transformation
3. Understand conditional P-values and Fisher's exact test for categorical data
4. Use the noncentral t distribution for exact power calculations
5. Determine sample sizes for planning clinical trials using R
6. Apply multiple comparison adjustments (Bonferroni, Benjamini-Hochberg)

Roadmap for Today

Topic	Key Idea	Textbook
P-values	Probability of data as extreme or more extreme than observed	Devore 9.4
P-value distribution	Uniform under H_0 ; concentrated near 0 under H_a	Devore 9.4
Fisher's exact test	Conditional P-values eliminate nuisance parameters	—
Power for the t test	Noncentral t distribution	Devore 9.2
Sample size determination	<code>power.t.test()</code> in R	Devore 9.2
Multiple testing	FWER (Bonferroni) and FDR (Benjamini-Hochberg)	—

Roadmap for Today

Topic	Key Idea	Textbook
P-values	Probability of data as extreme or more extreme than observed	Devore 9.4
P-value distribution	Uniform under H_0 ; concentrated near 0 under H_a	Devore 9.4
Fisher's exact test	Conditional P-values eliminate nuisance parameters	—
Power for the t test	Noncentral t distribution	Devore 9.2
Sample size determination	<code>power.t.test()</code> in R	Devore 9.2
Multiple testing	FWER (Bonferroni) and FDR (Benjamini-Hochberg)	—

Why This Matters

In medical research, P-values appear in virtually every published paper. Understanding what they mean — and what they do **not** mean — is essential for interpreting clinical evidence.

P-Values (Devore 9.4)

From Rejection Regions to P-Values

Recall the rejection region approach:

1. Fix α (e.g., 0.05)
2. Find critical value(s)
3. Reject H_0 if the test statistic falls in the rejection region

From Rejection Regions to P-Values

Recall the rejection region approach:

1. Fix α (e.g., 0.05)
2. Find critical value(s)
3. Reject H_0 if the test statistic falls in the rejection region

Limitation: The rejection region approach gives a binary decision (reject / fail to reject), but tells us nothing about **how strongly** the data contradict H_0 .

From Rejection Regions to P-Values

Recall the rejection region approach:

1. Fix α (e.g., 0.05)
2. Find critical value(s)
3. Reject H_0 if the test statistic falls in the rejection region

Limitation: The rejection region approach gives a binary decision (reject / fail to reject), but tells us nothing about **how strongly** the data contradict H_0 .

Two scenarios with the same decision:

- $\alpha = 0.05$, P-value = 0.049 $\Rightarrow$ barely reject H_0
- $\alpha = 0.05$, P-value = 0.00001 $\Rightarrow$ overwhelming evidence against H_0

The rejection region approach treats these identically!

Definition of the P-Value

! Definition (Devore 9.4)

The **P-value** is the probability, calculated assuming that the null hypothesis is true, of obtaining a value of the test statistic at least as contradictory to H_0 as the value calculated from the available sample.

Definition of the P-Value

! Definition (Devore 9.4)

The **P-value** is the probability, calculated assuming that the null hypothesis is true, of obtaining a value of the test statistic at least as contradictory to H_0 as the value calculated from the available sample.

Key points:

- The P-value **is** a probability
- It is calculated **assuming H_0 is true**
- It is a function of the **sample data** (and therefore a random variable!)
- We need to determine which values are “at least as contradictory” to H_0

Formal Definition of the P-Value

Let $T = T(X_1, \dots, X_n)$ be a test statistic and let t_{obs} be its observed value. For a test of $H_0 : \theta = \theta_0$ vs H_a :

! Definition (Formal)

The **P-value** is defined as:

$$\text{P-value} = \begin{cases} P_{\theta_0}(T \geq t_{obs}) & \text{upper-tailed test } (H_a : \theta > \theta_0) \\ P_{\theta_0}(T \leq t_{obs}) & \text{lower-tailed test } (H_a : \theta < \theta_0) \\ P_{\theta_0}(|T| \geq |t_{obs}|) & \text{two-tailed test } (H_a : \theta \neq \theta_0) \end{cases}$$

where $P_{\theta_0}(\cdot)$ denotes probability calculated under H_0 .

Formal Definition of the P-Value

Let $T = T(X_1, \dots, X_n)$ be a test statistic and let t_{obs} be its observed value. For a test of $H_0 : \theta = \theta_0$ vs H_a :

! Definition (Formal)

The **P-value** is defined as:

$$\text{P-value} = \begin{cases} P_{\theta_0}(T \geq t_{obs}) & \text{upper-tailed test } (H_a : \theta > \theta_0) \\ P_{\theta_0}(T \leq t_{obs}) & \text{lower-tailed test } (H_a : \theta < \theta_0) \\ P_{\theta_0}(|T| \geq |t_{obs}|) & \text{two-tailed test } (H_a : \theta \neq \theta_0) \end{cases}$$

where $P_{\theta_0}(\cdot)$ denotes probability calculated under H_0 .

Interpretation: The P-value measures how “extreme” or “surprising” the observed data would be if H_0 were true. Smaller P-values indicate data that is less compatible with H_0 .

The P-Value as a Tail Probability

Consider an upper-tailed test with test statistic T having CDF $F_T(t)$ under H_0 .

The P-Value as a Tail Probability

Consider an upper-tailed test with test statistic T having CDF $F_T(t)$ under H_0 .

The P-value is the “survival function” evaluated at t_{obs} :

$$\text{P-value} = P_{H_0}(T \geq t_{obs}) = 1 - F_T(t_{obs})$$

The P-Value as a Tail Probability

Consider an upper-tailed test with test statistic T having CDF $F_T(t)$ under H_0 .

The P-value is the “survival function” evaluated at t_{obs} :

$$\text{P-value} = P_{H_0}(T \geq t_{obs}) = 1 - F_T(t_{obs})$$

Key insight: The P-value is the probability of observing a test statistic *at least as extreme* as what we actually observed, computed under the assumption that H_0 is true.

The P-Value as a Tail Probability

Consider an upper-tailed test with test statistic T having CDF $F_T(t)$ under H_0 .

The P-value is the “survival function” evaluated at t_{obs} :

$$\text{P-value} = P_{H_0}(T \geq t_{obs}) = 1 - F_T(t_{obs})$$

Key insight: The P-value is the probability of observing a test statistic *at least as extreme* as what we actually observed, computed under the assumption that H_0 is true.

Why “at least as extreme”?

If t_{obs} is unlikely under H_0 , then values *even more extreme* than t_{obs} must also be unlikely. The P-value captures the total probability of all outcomes at least as unfavorable to H_0 .

The Probability Integral Transformation

A fundamental result from probability theory underlies the distribution of P-values.

! Theorem (Probability Integral Transformation)

Let X be a continuous random variable with CDF F_X . Then the random variable

$$U = F_X(X)$$

has a Uniform(0, 1) distribution.

The Probability Integral Transformation

A fundamental result from probability theory underlies the distribution of P-values.

! Theorem (Probability Integral Transformation)

Let X be a continuous random variable with CDF F_X . Then the random variable

$$U = F_X(X)$$

has a $\text{Uniform}(0, 1)$ distribution.

Proof: For $u \in (0, 1)$:

$$P(U \leq u) = P(F_X(X) \leq u) = P(X \leq F_X^{-1}(u)) = F_X(F_X^{-1}(u)) = u$$

The second equality uses the fact that F_X is monotone increasing (for continuous X). Since $P(U \leq u) = u$ for all $u \in (0, 1)$, we have $U \sim \text{Uniform}(0, 1)$. ■

Corollary: The Survival Function is Also Uniform

Corollary

If X is a continuous random variable with CDF F_X , then

$$V = 1 - F_X(X)$$

also has a Uniform(0, 1) distribution.

Corollary: The Survival Function is Also Uniform

Corollary

If X is a continuous random variable with CDF F_X , then

$$V = 1 - F_X(X)$$

also has a $\text{Uniform}(0, 1)$ distribution.

Proof: If $U \sim \text{Uniform}(0, 1)$, then $1 - U \sim \text{Uniform}(0, 1)$ as well.

This is easily verified: $P(1 - U \leq v) = P(U \geq 1 - v) = 1 - (1 - v) = v$. ■

Corollary: The Survival Function is Also Uniform

Corollary

If X is a continuous random variable with CDF F_X , then

$$V = 1 - F_X(X)$$

also has a $\text{Uniform}(0, 1)$ distribution.

Proof: If $U \sim \text{Uniform}(0, 1)$, then $1 - U \sim \text{Uniform}(0, 1)$ as well.

This is easily verified: $P(1 - U \leq v) = P(U \geq 1 - v) = 1 - (1 - v) = v$. ■

Why this matters: For an upper-tailed test, the P-value is $1 - F_T(T)$, which is exactly the survival function evaluated at the test statistic!

P-Values Are Uniform Under H_0

We can now state and prove the central result:

! Theorem (Distribution of the P-Value Under H_0)

Let T be a test statistic with continuous CDF F_T under H_0 . For an upper-tailed test, define the P-value as

$$P = 1 - F_T(T)$$

Then under H_0 , the P-value satisfies $P \sim \text{Uniform}(0, 1)$.

P-Values Are Uniform Under H_0

We can now state and prove the central result:

! Theorem (Distribution of the P-Value Under H_0)

Let T be a test statistic with continuous CDF F_T under H_0 . For an upper-tailed test, define the P-value as

$$P = 1 - F_T(T)$$

Then under H_0 , the P-value satisfies $P \sim \text{Uniform}(0, 1)$.

Proof: Under H_0 , the test statistic T has CDF F_T . By the probability integral transformation, $F_T(T) \sim \text{Uniform}(0, 1)$. By the corollary, $P = 1 - F_T(T) \sim \text{Uniform}(0, 1)$. ■

Consequence: P-Value Tests Maintain the Correct Size

This uniformity result has a crucial consequence for hypothesis testing.

Definition: Size of a Test

The **size** of a hypothesis test is the probability of rejecting H_0 when H_0 is true—i.e., the Type I error rate:

$$\text{Size} = P_{H_0}(\text{Reject } H_0) = P(\text{Type I Error})$$

A test has **exact size** α if this probability equals α exactly. A test has size **at most** α if this probability is $\leq \alpha$.

Consequence: P-Value Tests Maintain the Correct Size

This uniformity result has a crucial consequence for hypothesis testing.

Definition: Size of a Test

The **size** of a hypothesis test is the probability of rejecting H_0 when H_0 is true—i.e., the Type I error rate:

$$\text{Size} = P_{H_0}(\text{Reject } H_0) = P(\text{Type I Error})$$

A test has **exact size** α if this probability equals α exactly. A test has size **at most** α if this probability is $\leq \alpha$.

Theorem (Size of P-Value Tests)

If the P-value P has a $\text{Uniform}(0, 1)$ distribution under H_0 , then the test that rejects H_0 when $P \leq \alpha$ has **exact size** α .

Consequence: P-Value Tests Maintain the Correct Size

This uniformity result has a crucial consequence for hypothesis testing.

Definition: Size of a Test

The **size** of a hypothesis test is the probability of rejecting H_0 when H_0 is true—i.e., the Type I error rate:

$$\text{Size} = P_{H_0}(\text{Reject } H_0) = P(\text{Type I Error})$$

A test has **exact size** α if this probability equals α exactly. A test has size **at most** α if this probability is $\leq \alpha$.

Theorem (Size of P-Value Tests)

If the P-value P has a $\text{Uniform}(0, 1)$ distribution under H_0 , then the test that rejects H_0 when $P \leq \alpha$ has **exact size** α .

Proof: The Type I error rate is

$$P_{H_0}(\text{Reject } H_0) = P_{H_0}(P \leq \alpha) = \alpha$$

where the last equality follows because $P \sim \text{Uniform}(0, 1)$ under H_0 . ■

Consequence: P-Value Tests Maintain the Correct Size

This uniformity result has a crucial consequence for hypothesis testing.

Definition: Size of a Test

The **size** of a hypothesis test is the probability of rejecting H_0 when H_0 is true—i.e., the Type I error rate:

$$\text{Size} = P_{H_0}(\text{Reject } H_0) = P(\text{Type I Error})$$

A test has **exact size** α if this probability equals α exactly. A test has size **at most** α if this probability is $\leq \alpha$.

Theorem (Size of P-Value Tests)

If the P-value P has a $\text{Uniform}(0, 1)$ distribution under H_0 , then the test that rejects H_0 when $P \leq \alpha$ has **exact size** α .

Proof: The Type I error rate is

$$P_{H_0}(\text{Reject } H_0) = P_{H_0}(P \leq \alpha) = \alpha$$

where the last equality follows because $P \sim \text{Uniform}(0, 1)$ under H_0 . ■

Key insight: The P-value approach **automatically** guarantees that rejecting when P-value $\leq \alpha$ produces a test with Type I error rate exactly equal to α . This is not a coincidence—it's a direct consequence of the probability integral transformation!

Why This Matters: Equivalence of Approaches

Rejection region approach:

- Choose α , find critical value c such that $P_{H_0}(T \geq c) = \alpha$
- Reject H_0 if $T \geq c$

Why This Matters: Equivalence of Approaches

Rejection region approach:

- Choose α , find critical value c such that $P_{H_0}(T \geq c) = \alpha$
- Reject H_0 if $T \geq c$

P-value approach:

- Compute P-value = $P_{H_0}(T \geq t_{obs})$
- Reject H_0 if P-value $\leq \alpha$

Why This Matters: Equivalence of Approaches

Rejection region approach:

- Choose α , find critical value c such that $P_{H_0}(T \geq c) = \alpha$
- Reject H_0 if $T \geq c$

P-value approach:

- Compute P-value = $P_{H_0}(T \geq t_{obs})$
- Reject H_0 if P-value $\leq \alpha$

These are equivalent!

$$T \geq c \iff 1 - F_T(T) \leq 1 - F_T(c) = \alpha \iff \text{P-value} \leq \alpha$$

Why This Matters: Equivalence of Approaches

Rejection region approach:

- Choose α , find critical value c such that $P_{H_0}(T \geq c) = \alpha$
- Reject H_0 if $T \geq c$

P-value approach:

- Compute P-value = $P_{H_0}(T \geq t_{obs})$
- Reject H_0 if P-value $\leq \alpha$

These are equivalent!

$$T \geq c \iff 1 - F_T(T) \leq 1 - F_T(c) = \alpha \iff \text{P-value} \leq \alpha$$

The P-value encodes *all possible rejection decisions* for any choice of α :

- If P-value = 0.03: reject at $\alpha = 0.05$ but not at $\alpha = 0.01$
- If P-value = 0.008: reject at both $\alpha = 0.05$ and $\alpha = 0.01$

P-Values for Discrete Test Statistics

Technical Note

The probability integral transformation requires T to be **continuous**. For discrete test statistics (e.g., the binomial), the P-value may not be exactly Uniform $(0, 1)$.

P-Values for Discrete Test Statistics

Technical Note

The probability integral transformation requires T to be **continuous**. For discrete test statistics (e.g., the binomial), the P-value may not be exactly Uniform $(0, 1)$.

Example: For the exact binomial test with $H_0 : p = 0.5$ and $n = 10$, the P-value can only take finitely many values. The CDF of the P-value is a step function, not the identity.

P-Values for Discrete Test Statistics

Technical Note

The probability integral transformation requires T to be **continuous**. For discrete test statistics (e.g., the binomial), the P-value may not be exactly Uniform $(0, 1)$.

Example: For the exact binomial test with $H_0 : p = 0.5$ and $n = 10$, the P-value can only take finitely many values. The CDF of the P-value is a step function, not the identity.

Consequence: For discrete tests, $P(P \leq \alpha) \leq \alpha$, with equality only when α happens to coincide with one of the achievable P-values. The test is **conservative** (actual Type I error $<$ nominal α).

P-Values for Discrete Test Statistics

Technical Note

The probability integral transformation requires T to be **continuous**. For discrete test statistics (e.g., the binomial), the P-value may not be exactly Uniform $(0, 1)$.

Example: For the exact binomial test with $H_0 : p = 0.5$ and $n = 10$, the P-value can only take finitely many values. The CDF of the P-value is a step function, not the identity.

Consequence: For discrete tests, $P(P \leq \alpha) \leq \alpha$, with equality only when α happens to coincide with one of the achievable P-values. The test is **conservative** (actual Type I error $<$ nominal α).

Resolution: Randomized tests can achieve exact size α , but are rarely used in practice. Instead, we accept the slight conservatism.

Conditional P-Values and Fisher's Exact Test

What is Fisher's Exact Test?

Fisher's exact test is a statistical test used to determine whether there is a significant association between two categorical variables in a 2×2 contingency table.

What is Fisher's Exact Test?

Fisher's exact test is a statistical test used to determine whether there is a significant association between two categorical variables in a 2×2 contingency table.

Common applications in medical research:

- Is a genetic variant associated with disease risk?
- Does a treatment affect the probability of an adverse event?
- Is exposure to a risk factor associated with an outcome?
- Does a diagnostic test result predict disease status?

What is Fisher's Exact Test?

Fisher's exact test is a statistical test used to determine whether there is a significant association between two categorical variables in a 2×2 contingency table.

Common applications in medical research:

- Is a genetic variant associated with disease risk?
- Does a treatment affect the probability of an adverse event?
- Is exposure to a risk factor associated with an outcome?
- Does a diagnostic test result predict disease status?

Why "Exact"?

Unlike the chi-squared test (which relies on large-sample approximations), Fisher's exact test computes the **exact** P-value by enumerating all possible tables. This makes it particularly valuable for small samples or sparse tables where asymptotic approximations may be unreliable.

Fisher's Exact Test: The Setup

Consider data from a case-control or cross-sectional study:

	Outcome+	Outcome-	Total
Group A	a	b	a + b
Group B	c	d	c + d
Total	a + c	b + d	n

Fisher's Exact Test: The Setup

Consider data from a case-control or cross-sectional study:

	Outcome+	Outcome-	Total
Group A	a	b	a + b
Group B	c	d	c + d
Total	a + c	b + d	n

The question: Is the proportion with Outcome+ different between Group A and Group B?

Fisher's Exact Test: The Setup

Consider data from a case-control or cross-sectional study:

	Outcome+	Outcome-	Total
Group A	a	b	a + b
Group B	c	d	c + d
Total	a + c	b + d	n

The question: Is the proportion with Outcome+ different between Group A and Group B?

Equivalently: Are group membership and outcome **independent** or **associated**?

- H_0 : No association (independence): $P(\text{Outcome+} \mid \text{Group A}) = P(\text{Outcome+} \mid \text{Group B})$
- H_a : There is an association: $P(\text{Outcome+} \mid \text{Group A}) \neq P(\text{Outcome+} \mid \text{Group B})$

Fisher's Exact Test vs Chi-Squared Test

Feature	Fisher's Exact Test	Chi-Squared Test
P-value	Exact	Approximate (asymptotic)
Sample size	Valid for any n	Requires "large" n (expected counts ≥ 5)
Computation	Enumerates all tables	Simple formula
Sparse tables	Handles well	May be unreliable
One-sided tests	Natural	Awkward

Fisher's Exact Test vs Chi-Squared Test

Feature	Fisher's Exact Test	Chi-Squared Test
P-value	Exact	Approximate (asymptotic)
Sample size	Valid for any n	Requires "large" n (expected counts ≥ 5)
Computation	Enumerates all tables	Simple formula
Sparse tables	Handles well	May be unreliable
One-sided tests	Natural	Awkward

 When to Use Fisher's Exact Test

- Any expected cell count < 5
- Total sample size < 20 -30
- When exact P-values are required (regulatory submissions)
- One-sided hypotheses about association direction

For large samples with adequate expected counts, chi-squared and Fisher's give nearly identical results.

The Theoretical Challenge

To understand Fisher's exact test deeply, we need to address a fundamental question:

The Theoretical Challenge

To understand Fisher's exact test deeply, we need to address a fundamental question:

How do we compute a P-value when the null hypothesis doesn't fully specify the distribution?

The Theoretical Challenge

To understand Fisher's exact test deeply, we need to address a fundamental question:

How do we compute a P-value when the null hypothesis doesn't fully specify the distribution?

Testing H_0 : "Group and outcome are independent" does **not** tell us:

- What proportion of subjects are in Group A
- What proportion have the outcome

These **nuisance parameters** affect the distribution of any test statistic, yet we don't know their values!

The Theoretical Challenge

To understand Fisher's exact test deeply, we need to address a fundamental question:

How do we compute a P-value when the null hypothesis doesn't fully specify the distribution?

Testing H_0 : "Group and outcome are independent" does **not** tell us:

- What proportion of subjects are in Group A
- What proportion have the outcome

These **nuisance parameters** affect the distribution of any test statistic, yet we don't know their values!

Fisher's solution: Condition on the sufficient statistics for the nuisance parameters, eliminating them from the problem entirely.

The Problem of Nuisance Parameters

In many testing problems, the null hypothesis specifies a relationship but leaves some parameters unspecified.

The Problem of Nuisance Parameters

In many testing problems, the null hypothesis specifies a relationship but leaves some parameters unspecified.

Example: Testing independence in a 2×2 table

	Disease+	Disease−	Total
Exposed	a	b	a + b
Not Exposed	c	d	c + d
Total	a + c	b + d	n

The Problem of Nuisance Parameters

In many testing problems, the null hypothesis specifies a relationship but leaves some parameters unspecified.

Example: Testing independence in a 2×2 table

	Disease+	Disease−	Total
Exposed	a	b	a + b
Not Exposed	c	d	c + d
Total	a + c	b + d	n

Hypotheses:

- H_0 : Exposure and disease are independent (no association)
- H_a : Exposure and disease are associated

The Problem of Nuisance Parameters

In many testing problems, the null hypothesis specifies a relationship but leaves some parameters unspecified.

Example: Testing independence in a 2×2 table

	Disease+	Disease−	Total
Exposed	a	b	a + b
Not Exposed	c	d	c + d
Total	a + c	b + d	n

Hypotheses:

- H_0 : Exposure and disease are independent (no association)
- H_a : Exposure and disease are associated

The problem: Under H_0 , the marginal probabilities $p_1 = P(\text{Exposed})$ and $p_2 = P(\text{Disease+})$ are **nuisance parameters**—unknown quantities not of direct interest.

Why Nuisance Parameters Complicate P-Values

To compute a P-value, we need to calculate $P_{H_0}(T \geq t_{obs})$.

Why Nuisance Parameters Complicate P-Values

To compute a P-value, we need to calculate $P_{H_0}(T \geq t_{obs})$.

But under $H_0 : p_{11} = p_1 \cdot p_2$ (independence), the distribution of any test statistic depends on the unknown values of p_1 and p_2 !

Why Nuisance Parameters Complicate P-Values

To compute a P-value, we need to calculate $P_{H_0}(T \geq t_{obs})$.

But under $H_0 : p_{11} = p_1 \cdot p_2$ (independence), the distribution of any test statistic depends on the unknown values of p_1 and p_2 !

Options:

1. **Estimate** the nuisance parameters and use asymptotic approximations (chi-squared test)
2. **Eliminate** the nuisance parameters by conditioning on sufficient statistics

Why Nuisance Parameters Complicate P-Values

To compute a P-value, we need to calculate $P_{H_0}(T \geq t_{obs})$.

But under $H_0 : p_{11} = p_1 \cdot p_2$ (independence), the distribution of any test statistic depends on the unknown values of p_1 and p_2 !

Options:

1. **Estimate** the nuisance parameters and use asymptotic approximations (chi-squared test)
2. **Eliminate** the nuisance parameters by conditioning on sufficient statistics

Fisher's exact test takes approach (2), which yields an **exact** P-value that is valid for any sample size.

The Conditionality Principle

! Conditionality Principle (Informal)

If a statistic S is **ancillary** (its distribution doesn't depend on the parameter of interest), then inference should be conducted **conditional** on the observed value of S .

The Conditionality Principle

! Conditionality Principle (Informal)

If a statistic S is **ancillary** (its distribution doesn't depend on the parameter of interest), then inference should be conducted **conditional** on the observed value of S .

Intuition: Ancillary statistics carry no information about the parameter of interest, but they do affect the “precision” of the experiment. Conditioning on them focuses inference on the relevant variation.

The Conditionality Principle

! Conditionality Principle (Informal)

If a statistic S is **ancillary** (its distribution doesn't depend on the parameter of interest), then inference should be conducted **conditional** on the observed value of S .

Intuition: Ancillary statistics carry no information about the parameter of interest, but they do affect the “precision” of the experiment. Conditioning on them focuses inference on the relevant variation.

For nuisance parameters: If S is a sufficient statistic for the nuisance parameter η , then conditioning on S eliminates the dependence on η .

Fisher's Exact Test: Setup

Consider a 2×2 table with fixed total n . Let:

- X = number in cell (1,1) (Exposed and Disease+)
- Row totals $r_1 = a + b$ and $r_2 = c + d$ (fixed by design or conditioning)
- Column totals $c_1 = a + c$ and $c_2 = b + d$

Fisher's Exact Test: Setup

Consider a 2×2 table with fixed total n . Let:

- X = number in cell (1,1) (Exposed and Disease+)
- Row totals $r_1 = a + b$ and $r_2 = c + d$ (fixed by design or conditioning)
- Column totals $c_1 = a + c$ and $c_2 = b + d$

Under H_0 (independence) with both margins fixed, the cell count X follows a **hypergeometric distribution**:

$$P(X = x \mid r_1, r_2, c_1, c_2) = \frac{\binom{c_1}{x} \binom{c_2}{r_1 - x}}{\binom{n}{r_1}}$$

Fisher's Exact Test: Setup

Consider a 2×2 table with fixed total n . Let:

- X = number in cell (1,1) (Exposed and Disease+)
- Row totals $r_1 = a + b$ and $r_2 = c + d$ (fixed by design or conditioning)
- Column totals $c_1 = a + c$ and $c_2 = b + d$

Under H_0 (independence) with both margins fixed, the cell count X follows a **hypergeometric distribution**:

$$P(X = x \mid r_1, r_2, c_1, c_2) = \frac{\binom{c_1}{x} \binom{c_2}{r_1 - x}}{\binom{n}{r_1}}$$

Key observation: This distribution has **no unknown parameters!** The nuisance parameters have been eliminated.

Why Condition on the Margins?

Theorem (Sufficiency of Margins)

Under H_0 (independence), the row and column marginal totals (r_1, c_1) are **jointly sufficient** for the nuisance parameters (p_1, p_2) .

Why Condition on the Margins?

Theorem (Sufficiency of Margins)

Under H_0 (independence), the row and column marginal totals (r_1, c_1) are **jointly sufficient** for the nuisance parameters (p_1, p_2) .

Proof sketch: Under independence, the joint probability of a 2×2 table is:

$$P(a, b, c, d) = \frac{n!}{a!b!c!d!} p_1^{r_1} (1 - p_1)^{r_2} p_2^{c_1} (1 - p_2)^{c_2}$$

Why Condition on the Margins?

Theorem (Sufficiency of Margins)

Under H_0 (independence), the row and column marginal totals (r_1, c_1) are **jointly sufficient** for the nuisance parameters (p_1, p_2) .

Proof sketch: Under independence, the joint probability of a 2×2 table is:

$$P(a, b, c, d) = \frac{n!}{a!b!c!d!} p_1^{r_1} (1 - p_1)^{r_2} p_2^{c_1} (1 - p_2)^{c_2}$$

By the factorization theorem, since the likelihood factors into a term depending on (p_1, p_2) only through (r_1, c_1) , the margins are sufficient for the nuisance parameters.

The Conditional Distribution Eliminates Nuisance Parameters

Conditioning on the sufficient statistic (r_1, c_1) :

$$P(X = x \mid r_1, c_1) = \frac{P(X = x, r_1, c_1)}{P(r_1, c_1)}$$

The Conditional Distribution Eliminates Nuisance Parameters

Conditioning on the sufficient statistic (r_1, c_1) :

$$P(X = x \mid r_1, c_1) = \frac{P(X = x, r_1, c_1)}{P(r_1, c_1)}$$

After simplification (the nuisance parameters cancel!):

$$P(X = x \mid r_1, c_1) = \frac{\binom{c_1}{x} \binom{n-c_1}{r_1-x}}{\binom{n}{r_1}}$$

The Conditional Distribution Eliminates Nuisance Parameters

Conditioning on the sufficient statistic (r_1, c_1) :

$$P(X = x \mid r_1, c_1) = \frac{P(X = x, r_1, c_1)}{P(r_1, c_1)}$$

After simplification (the nuisance parameters cancel!):

$$P(X = x \mid r_1, c_1) = \frac{\binom{c_1}{x} \binom{n-c_1}{r_1-x}}{\binom{n}{r_1}}$$

This is the hypergeometric distribution with parameters n , r_1 , and c_1 —all **known** from the observed data!

The Conditional Distribution Eliminates Nuisance Parameters

Conditioning on the sufficient statistic (r_1, c_1) :

$$P(X = x \mid r_1, c_1) = \frac{P(X = x, r_1, c_1)}{P(r_1, c_1)}$$

After simplification (the nuisance parameters cancel!):

$$P(X = x \mid r_1, c_1) = \frac{\binom{c_1}{x} \binom{n-c_1}{r_1-x}}{\binom{n}{r_1}}$$

This is the hypergeometric distribution with parameters n , r_1 , and c_1 —all **known** from the observed data!

Key Insight

By conditioning on the margins, we transform a problem with nuisance parameters into one where the null distribution is completely specified.

Fisher's Exact Test: The Conditional P-Value

The **conditional P-value** is computed using the hypergeometric distribution:

$$\text{P-value} = P(X \geq x_{obs} \mid r_1, c_1) = \sum_{x=x_{obs}}^{\min(r_1, c_1)} \frac{\binom{c_1}{x} \binom{n-c_1}{r_1-x}}{\binom{n}{r_1}}$$

Fisher's Exact Test: The Conditional P-Value

The **conditional P-value** is computed using the hypergeometric distribution:

$$\text{P-value} = P(X \geq x_{obs} \mid r_1, c_1) = \sum_{x=x_{obs}}^{\min(r_1, c_1)} \frac{\binom{c_1}{x} \binom{n-c_1}{r_1-x}}{\binom{n}{r_1}}$$

(For a two-sided test, we sum probabilities for outcomes as extreme or more extreme in either direction.)

Fisher's Exact Test: The Conditional P-Value

The **conditional P-value** is computed using the hypergeometric distribution:

$$\text{P-value} = P(X \geq x_{obs} \mid r_1, c_1) = \sum_{x=x_{obs}}^{\min(r_1, c_1)} \frac{\binom{c_1}{x} \binom{n-c_1}{r_1-x}}{\binom{n}{r_1}}$$

(For a two-sided test, we sum probabilities for outcomes as extreme or more extreme in either direction.)

Properties:

- **Exact:** Valid for any sample size, including small samples
- **Free of nuisance parameters:** No need to estimate marginal probabilities
- **Discrete:** P-value can only take finitely many values (test is slightly conservative)

Example: Fisher's Exact Test in R

Setting: A clinical trial investigates whether a new treatment reduces adverse events.

```
1 # Observed 2x2 table
2 #
3 # Treatment      2      18      (20)
4 # Placebo        8      12      (20)
5 #                (10)     (30)     (40)
6
7 adverse_table <- matrix(c(2, 8, 18, 12), nrow = 2, byrow = FALSE,
8                             dimnames = list(Group = c("Treatment", "Placebo"),
9                                                  Outcome = c("Adverse", "None")))
10 adverse_table
```

Group	Outcome	
	Adverse	None
Treatment	2	18
Placebo	8	12

Example: Fisher's Exact Test Output

```
1 # Fisher's exact test
2 fisher_result <- fisher.test(adverse_table)
3 fisher_result
```

Fisher's Exact Test for Count Data

```
data:  adverse_table
p-value = 0.06483
alternative hypothesis: true odds ratio is not equal to 1
95 percent confidence interval:
 0.01546086 1.08434068
sample estimates:
odds ratio
 0.1743672
```


Example: Fisher's Exact Test Output

```
1 # Fisher's exact test
2 fisher_result <- fisher.test(adverse_table)
3 fisher_result
```

Fisher's Exact Test for Count Data

```
data: adverse_table
p-value = 0.06483
alternative hypothesis: true odds ratio is not equal to 1
95 percent confidence interval:
 0.01546086 1.08434068
sample estimates:
odds ratio
 0.1743672
```

```
1 cat("P-value:", round(fisher_result$p.value, 4), "\n")
```

P-value: 0.0648

```
1 cat("Odds ratio estimate:", round(fisher_result$estimate, 3), "\n")
```

Odds ratio estimate: 0.174

```
1 cat("95% CI for odds ratio:", round(fisher_result$conf.int[1], 3), "to",
2     round(fisher_result$conf.int[2], 3), "\n")
```

95% CI for odds ratio: 0.015 to 1.084




Visualizing the Hypergeometric Distribution

► Code

Hypergeometric Distribution Under H_0 (Conditional on Margins)

$n = 40$, $r_1 = 20$, $c_1 = 10$; P-value = 0.0648



Why Conditioning? A Deeper Look

The philosophical justification:

1. **Sufficiency Principle:** Inference should depend on the data only through sufficient statistics
2. **Conditionality Principle:** Inference should condition on ancillary statistics

Why Conditioning? A Deeper Look

The philosophical justification:

1. **Sufficiency Principle:** Inference should depend on the data only through sufficient statistics
2. **Conditionality Principle:** Inference should condition on ancillary statistics

When combined, these principles suggest conditioning on sufficient statistics for nuisance parameters.

Why Conditioning? A Deeper Look

The philosophical justification:

1. **Sufficiency Principle:** Inference should depend on the data only through sufficient statistics
2. **Conditionality Principle:** Inference should condition on ancillary statistics

When combined, these principles suggest conditioning on sufficient statistics for nuisance parameters.

The practical justification:

- Conditioning produces a test with **exact** size α
- No large-sample approximations needed
- Valid even for sparse tables or small samples

Why Conditioning? A Deeper Look

The philosophical justification:

1. **Sufficiency Principle:** Inference should depend on the data only through sufficient statistics
2. **Conditionality Principle:** Inference should condition on ancillary statistics

When combined, these principles suggest conditioning on sufficient statistics for nuisance parameters.

The practical justification:

- Conditioning produces a test with **exact** size α
- No large-sample approximations needed
- Valid even for sparse tables or small samples

Caveat

Conditioning can reduce power compared to unconditional tests. This is the price of exactness. For large samples, the chi-squared test (which estimates nuisance parameters) may be more powerful.

Connection to Randomization Tests

Fisher's exact test can also be viewed as a **randomization test**:

Connection to Randomization Tests

Fisher's exact test can also be viewed as a **randomization test**:

Setup: Fix the margins (total exposed, total diseased). Under H_0 , each possible table consistent with these margins is equally likely given the hypergeometric distribution.

Connection to Randomization Tests

Fisher's exact test can also be viewed as a **randomization test**:

Setup: Fix the margins (total exposed, total diseased). Under H_0 , each possible table consistent with these margins is equally likely given the hypergeometric distribution.

The P-value is the proportion of possible tables that show an association at least as extreme as observed.

Connection to Randomization Tests

Fisher's exact test can also be viewed as a **randomization test**:

Setup: Fix the margins (total exposed, total diseased). Under H_0 , each possible table consistent with these margins is equally likely given the hypergeometric distribution.

The P-value is the proportion of possible tables that show an association at least as extreme as observed.

This connects to permutation testing: Under H_0 , "Outcome" labels are exchangeable among subjects. The conditional P-value equals the proportion of permutations yielding a test statistic as extreme as observed.

```
1 # Verify: Fisher's P-value via permutation simulation (one-sided version)
2 set.seed(551)
3 n_perm <- 10000
4
5 # Reconstruct individual-level data matching adverse_table:
6 # Treatment: 2 adverse, 18 none; Placebo: 8 adverse, 12 none
7 treatment <- c(rep(1, 20), rep(0, 20)) # 20 treatment, 20 placebo
8 adverse <- c(rep(1, 2), rep(0, 18),      # Treatment: 2 adverse, 18 not
9             rep(1, 8), rep(0, 12))      # Placebo: 8 adverse, 12 not
10
11 # Observed statistic: treatment patients with adverse events
12 obs_stat <- sum(treatment * adverse) # = 2
13 cat("Observed statistic (treatment + adverse):", obs_stat, "\n")
```

Observed statistic (treatment + adverse): 2

Fisher's Exact Test: Complete R Workflow

Clinical scenario: A pilot study investigates whether a genetic variant is associated with response to a targeted cancer therapy.

```
1 # Create the contingency table from raw data
2 patients <- tibble(
3   patient_id = 1:45,
4   variant = c(rep("Present", 18), rep("Absent", 27)),
5   response = c(
6     # Variant present: 12 responders, 6 non-responders
7     rep("Responder", 12), rep("Non-responder", 6),
8     # Variant absent: 8 responders, 19 non-responders
9     rep("Responder", 8), rep("Non-responder", 19)
10  )
11 )
12
13 # Cross-tabulation
14 response_table <- table(patients$variant, patients$response)
15 response_table
```

	Non-responder	Responder
Absent	19	8
Present	6	12

Fisher's Exact Test: Analysis

```
1 # Fisher's exact test (two-sided by default)
2 fisher_result <- fisher.test(response_table)
3
4 # Display full output
5 fisher_result
```

Fisher's Exact Test for Count Data

```
data: response_table
p-value = 0.03071
alternative hypothesis: true odds ratio is not equal to 1
95 percent confidence interval:
 1.121584 20.940769
sample estimates:
odds ratio
 4.570771
```

Fisher's Exact Test: Extracting Results

```
1 # Extract components for reporting
2 cat("=== Fisher's Exact Test Results ===\n\n")
```

=== Fisher's Exact Test Results ===

```
1 cat("P-value:", format(fisher_result$p.value, digits = 4), "\n")
```

P-value: 0.03071

```
1 cat("Odds Ratio:", round(fisher_result$estimate, 2), "\n")
```

Odds Ratio: 4.57

```
1 cat("95% CI for OR: [", round(fisher_result$conf.int[1], 2), ", ",
2     round(fisher_result$conf.int[2], 2), "]\n\n")
```

95% CI for OR: [1.12 , 20.94]

```
1 # Decision at alpha = 0.05
2 alpha <- 0.05
3 if (fisher_result$p.value <= alpha) {
4   cat("Decision: Reject H0 at alpha =", alpha, "\n")
5   cat("Conclusion: There is a statistically significant association\n")
6   cat("between the genetic variant and treatment response.\n")
7 } else {
8   cat("Decision: Fail to reject H0 at alpha =", alpha, "\n")
9   cat("Conclusion: Insufficient evidence of association.\n")
10 }
```

Fisher's Exact Test: One-Sided Alternatives

```
1 # If we have a directional hypothesis (variant improves response)
2 fisher_greater <- fisher.test(response_table, alternative = "greater")
3 fisher_less <- fisher.test(response_table, alternative = "less")
4
5 cat("Two-sided P-value:", round(fisher_result$p.value, 4), "\n")
```

Two-sided P-value: 0.0307

```
1 cat("One-sided (OR > 1):", round(fisher_greater$p.value, 4), "\n")
```

One-sided (OR > 1): 0.0157

```
1 cat("One-sided (OR < 1):", round(fisher_less$p.value, 4), "\n")
```

One-sided (OR < 1): 0.9973

Fisher's Exact Test: One-Sided Alternatives

```
1 # If we have a directional hypothesis (variant improves response)
2 fisher_greater <- fisher.test(response_table, alternative = "greater")
3 fisher_less <- fisher.test(response_table, alternative = "less")
4
5 cat("Two-sided P-value:", round(fisher_result$p.value, 4), "\n")
```

Two-sided P-value: 0.0307

```
1 cat("One-sided (OR > 1):", round(fisher_greater$p.value, 4), "\n")
```

One-sided (OR > 1): 0.0157

```
1 cat("One-sided (OR < 1):", round(fisher_less$p.value, 4), "\n")
```

One-sided (OR < 1): 0.9973

Choosing One-Sided vs Two-Sided

Use a **one-sided** test only when you have a strong *a priori* hypothesis about the direction of association, specified **before** seeing the data. In most exploratory studies, two-sided tests are more appropriate.

Comparing Fisher's Exact Test to Chi-Squared Test

```
1 # Chi-squared test (asymptotic approximation)
2 chisq_result <- chisq.test(response_table)
3 chisq_result
```

Pearson's Chi-squared test with Yates' continuity correction

data: response_table

X-squared = 4.5938, df = 1, p-value = 0.03209

Comparing Fisher's Exact Test to Chi-Squared Test

```
1 # Chi-squared test (asymptotic approximation)
2 chisq_result <- chisq.test(response_table)
3 chisq_result
```

Pearson's Chi-squared test with Yates' continuity correction

```
data: response_table
X-squared = 4.5938, df = 1, p-value = 0.03209
```

```
1 # Compare P-values
2 cat("\nComparison of P-values:\n")
```

Comparison of P-values:

```
1 cat("Fisher's exact test:", round(fisher_result$p.value, 4), "\n")
```

Fisher's exact test: 0.0307

```
1 cat("Chi-squared test: ", round(chisq_result$p.value, 4), "\n")
```

Chi-squared test: 0.0321

```
1 cat("\nExpected counts (chi-squared assumes these are >= 5):\n")
```

Expected counts (chi-squared assumes these are ≥ 5):

```
1 round(chisq_result$expected, 1)
```

Comparing Fisher's Exact Test to Chi-Squared Test

```
1 # Chi-squared test (asymptotic approximation)
2 chisq_result <- chisq.test(response_table)
3 chisq_result
```

Pearson's Chi-squared test with Yates' continuity correction

```
data: response_table
X-squared = 4.5938, df = 1, p-value = 0.03209
```

```
1 # Compare P-values
2 cat("\nComparison of P-values:\n")
```

Comparison of P-values:

```
1 cat("Fisher's exact test:", round(fisher_result$p.value, 4), "\n")
```

Fisher's exact test: 0.0307

```
1 cat("Chi-squared test: ", round(chisq_result$p.value, 4), "\n")
```

Chi-squared test: 0.0321

```
1 cat("\nExpected counts (chi-squared assumes these are >= 5):\n")
```

Expected counts (chi-squared assumes these are ≥ 5):

```
1 round(chisq_result$expected, 1)
```

Decision Rule Based on the P-Value

! P-Value Decision Rule (Devore 9.4)

Select a significance level α . Then:

Reject H_0 if P-value $\leq \alpha$; Do not reject H_0 if P-value $> \alpha$

Decision Rule Based on the P-Value

! P-Value Decision Rule (Devore 9.4)

Select a significance level α . Then:

Reject H_0 if P-value $\leq \alpha$; Do not reject H_0 if P-value $> \alpha$

Equivalence to the rejection region method:

- If the test statistic falls in the rejection region, then P-value $\leq \alpha$
- If the test statistic falls outside the rejection region, then P-value $> \alpha$

Decision Rule Based on the P-Value

! P-Value Decision Rule (Devore 9.4)

Select a significance level α . Then:

Reject H_0 if P-value $\leq \alpha$; Do not reject H_0 if P-value $> \alpha$

Equivalence to the rejection region method:

- If the test statistic falls in the rejection region, then P-value $\leq \alpha$
- If the test statistic falls outside the rejection region, then P-value $> \alpha$

Advantage: The P-value encodes more information than a binary reject/fail-to-reject decision. It conveys the **strength** of evidence against H_0 .

P-Values for z Tests

For a z test (test statistic Z with $Z \sim N(0, 1)$ under H_0):

$$\text{P-value} = \begin{cases} 1 - \Phi(z) & \text{for an upper-tailed test } (H_a : \mu > \mu_0) \\ \Phi(z) & \text{for a lower-tailed test } (H_a : \mu < \mu_0) \\ 2[1 - \Phi(|z|)] & \text{for a two-tailed test } (H_a : \mu \neq \mu_0) \end{cases}$$

where $\Phi(z)$ is the standard normal CDF and z is the observed test statistic.

P-Values for z Tests

For a z test (test statistic Z with $Z \sim N(0, 1)$ under H_0):

$$\text{P-value} = \begin{cases} 1 - \Phi(z) & \text{for an upper-tailed test } (H_a : \mu > \mu_0) \\ \Phi(z) & \text{for a lower-tailed test } (H_a : \mu < \mu_0) \\ 2[1 - \Phi(|z|)] & \text{for a two-tailed test } (H_a : \mu \neq \mu_0) \end{cases}$$

where $\Phi(z)$ is the standard normal CDF and z is the observed test statistic.

```
1 # Quick R reminders for P-value computation
2 z_obs <- 2.15
3 cat("Upper-tailed P-value:", 1 - pnorm(z_obs), "\n")
```

Upper-tailed P-value: 0.01577761

```
1 cat("Lower-tailed P-value:", pnorm(z_obs), "\n")
```


Lower-tailed P-value: 0.9842224

```
1 cat("Two-tailed P-value: ", 2 * (1 - pnorm(abs(z_obs))), "\n")
```

Two-tailed P-value: 0.03155521

Visualizing P-Values for z Tests

► Code



P-Values for t Tests

For the one-sample t test with $n - 1$ degrees of freedom, the P-value formulas are analogous:

$$\text{P-value} = \begin{cases} P(T_{n-1} \geq t) & \text{upper-tailed} \\ P(T_{n-1} \leq t) & \text{lower-tailed} \\ 2 \cdot P(T_{n-1} \geq |t|) & \text{two-tailed} \end{cases}$$

P-Values for t Tests

For the one-sample t test with $n - 1$ degrees of freedom, the P-value formulas are analogous:

$$\text{P-value} = \begin{cases} P(T_{n-1} \geq t) & \text{upper-tailed} \\ P(T_{n-1} \leq t) & \text{lower-tailed} \\ 2 \cdot P(T_{n-1} \geq |t|) & \text{two-tailed} \end{cases}$$

```
1 # P-values for t tests in R
2 t_obs <- 2.15
3 df <- 14
4
5 cat("Upper-tailed P-value:", pt(t_obs, df = df, lower.tail = FALSE), "\n")
```

Upper-tailed P-value: 0.0247585

```
1 cat("Lower-tailed P-value:", pt(t_obs, df = df, lower.tail = TRUE), "\n")
```

Lower-tailed P-value: 0.9752415

```
1 cat("Two-tailed P-value: ", 2 * pt(abs(t_obs), df = df, lower.tail = FALSE), "\n")
```

Two-tailed P-value: 0.04951701

P-Values for t Tests

For the one-sample t test with $n - 1$ degrees of freedom, the P-value formulas are analogous:

$$\text{P-value} = \begin{cases} P(T_{n-1} \geq t) & \text{upper-tailed} \\ P(T_{n-1} \leq t) & \text{lower-tailed} \\ 2 \cdot P(T_{n-1} \geq |t|) & \text{two-tailed} \end{cases}$$

```
1 # P-values for t tests in R
2 t_obs <- 2.15
3 df <- 14
4
5 cat("Upper-tailed P-value:", pt(t_obs, df = df, lower.tail = FALSE), "\n")
```

Upper-tailed P-value: 0.0247585

```
1 cat("Lower-tailed P-value:", pt(t_obs, df = df, lower.tail = TRUE), "\n")
```

Lower-tailed P-value: 0.9752415

```
1 cat("Two-tailed P-value: ", 2 * pt(abs(t_obs), df = df, lower.tail = FALSE), "\n")
```

Two-tailed P-value: 0.04951701

R Functions

`t.test()` and `prop.test()` automatically compute P-values. Manual computation is needed only for understanding and for non-standard situations.



Medical Examples: Computing P-Values

Example: Hemoglobin A1c Levels in Diabetic Patients

Setting: A diabetes clinic is evaluating whether a new management protocol changes mean hemoglobin A1c (HbA1c) levels. The standard target is 7.0%. A sample of $n = 20$ patients on the new protocol gives $\bar{x} = 7.42\%$ and $s = 0.95\%$.

Question: Is there evidence that the new protocol produces different mean HbA1c than the target?

Example: Hemoglobin A1c Levels in Diabetic Patients

Setting: A diabetes clinic is evaluating whether a new management protocol changes mean hemoglobin A1c (HbA1c) levels. The standard target is 7.0%. A sample of $n = 20$ patients on the new protocol gives $\bar{x} = 7.42\%$ and $s = 0.95\%$.

Question: Is there evidence that the new protocol produces different mean HbA1c than the target?

```
1 # Data summary
2 n <- 20
3 xbar <- 7.42
4 s <- 0.95
5 mu0 <- 7.0
6
7 # Test statistic
8 t_stat <- (xbar - mu0) / (s / sqrt(n))
9 cat("Test statistic: t =", round(t_stat, 3), "\n")
```

Test statistic: t = 1.977

```
1 cat("Degrees of freedom:", n - 1, "\n")
```

Degrees of freedom: 19

```
1 # P-value (two-tailed)
2 p_value <- 2 * pt(abs(t_stat), df = n - 1, lower.tail = FALSE)
3 cat("P-value (two-tailed):", round(p_value, 4), "\n")
```

P-value (two-tailed): 0.0627




HbA1c Example: Visualizing the P-Value

► Code

Two-Tailed t Test: HbA1c Levels (t_{19} distribution)

$H_0: \mu = 7\%$ vs $H_a: \mu \neq 7\%$



HbA1c Example: Complete Analysis with `t.test()`

```
1 # Simulate individual patient data consistent with our summary statistics
2 set.seed(2026)
3 hba1c <- rnorm(20, mean = 7.42, sd = 0.95)
4 # Adjust to match desired summary statistics closely
5 hba1c <- (hba1c - mean(hba1c)) / sd(hba1c) * 0.95 + 7.42
6
7 # Full t.test output
8 hba1c_test <- t.test(hba1c, mu = 7.0, alternative = "two.sided")
9 hba1c_test
```

One Sample t-test

```
data:  hba1c
t = 1.9772, df = 19, p-value = 0.06272
alternative hypothesis: true mean is not equal to 7
95 percent confidence interval:
 6.975386 7.864614
sample estimates:
mean of x
 7.42
```


HbA1c Example: Complete Analysis with `t.test()`

```
1 # Simulate individual patient data consistent with our summary statistics
2 set.seed(2026)
3 hba1c <- rnorm(20, mean = 7.42, sd = 0.95)
4 # Adjust to match desired summary statistics closely
5 hba1c <- (hba1c - mean(hba1c)) / sd(hba1c) * 0.95 + 7.42
6
7 # Full t.test output
8 hba1c_test <- t.test(hba1c, mu = 7.0, alternative = "two.sided")
9 hba1c_test
```

One Sample t-test

```
data: hba1c
t = 1.9772, df = 19, p-value = 0.06272
alternative hypothesis: true mean is not equal to 7
95 percent confidence interval:
 6.975386 7.864614
sample estimates:
mean of x
 7.42
```

Interpretation: At $\alpha = 0.05$, the P-value of 0.0627 is less than 0.05, so we reject H_0 . There is statistically significant evidence that the mean HbA1c under the new protocol differs from 7.0%.



Example: Serum Transferrin Receptor in Pregnancy (Devore 9.4, Exercise 55)

Setting: The average serum transferrin receptor concentration for all women is known to be 5.63 mg/L. A study reported P-value > 0.10 for testing $H_0 : \mu = 5.63$ vs $H_a : \mu \neq 5.63$ based on $n = 176$ pregnant women.

Example: Serum Transferrin Receptor in Pregnancy (Devore 9.4, Exercise 55)

Setting: The average serum transferrin receptor concentration for all women is known to be 5.63 mg/L. A study reported P-value > 0.10 for testing $H_0 : \mu = 5.63$ vs $H_a : \mu \neq 5.63$ based on $n = 176$ pregnant women.

Question: Using $\alpha = 0.01$, what would you conclude?

Example: Serum Transferrin Receptor in Pregnancy (Devore 9.4, Exercise 55)

Setting: The average serum transferrin receptor concentration for all women is known to be 5.63 mg/L. A study reported P-value > 0.10 for testing $H_0 : \mu = 5.63$ vs $H_a : \mu \neq 5.63$ based on $n = 176$ pregnant women.

Question: Using $\alpha = 0.01$, what would you conclude?

Conclusion: There is insufficient evidence at the 0.01 significance level to conclude that the mean serum transferrin receptor concentration for pregnant women differs from the value of 5.63 mg/L for all women.

Example: Post-Surgical Recovery Time

Setting: The standard mean recovery time after a particular surgery is 14 days. A surgeon has adopted a new minimally invasive technique and records recovery times for 15 patients.

```
1 recovery <- c(11.2, 13.5, 10.8, 12.1, 14.3, 9.7, 11.9, 12.8,  
2              10.5, 13.1, 11.4, 12.6, 10.2, 13.8, 11.1)  
3  
4 # One-sided test:  $H_0: \mu = 14$  vs  $H_a: \mu < 14$   
5 recovery_test <- t.test(recovery, mu = 14, alternative = "less")  
6 recovery_test
```

One Sample t-test

```
data:  recovery  
t = -5.7665, df = 14, p-value = 2.441e-05  
alternative hypothesis: true mean is less than 14  
95 percent confidence interval:  
    -Inf 12.56457  
sample estimates:  
mean of x  
11.93333
```

Example: Post-Surgical Recovery Time

Setting: The standard mean recovery time after a particular surgery is 14 days. A surgeon has adopted a new minimally invasive technique and records recovery times for 15 patients.

```
1 recovery <- c(11.2, 13.5, 10.8, 12.1, 14.3, 9.7, 11.9, 12.8,  
2              10.5, 13.1, 11.4, 12.6, 10.2, 13.8, 11.1)  
3  
4 # One-sided test:  $H_0: \mu = 14$  vs  $H_a: \mu < 14$   
5 recovery_test <- t.test(recovery, mu = 14, alternative = "less")  
6 recovery_test
```

One Sample t-test

```
data:  recovery  
t = -5.7665, df = 14, p-value = 2.441e-05  
alternative hypothesis: true mean is less than 14  
95 percent confidence interval:  
    -Inf 12.56457  
sample estimates:  
mean of x  
11.93333
```

```
1 cat("P-value =", format(recovery_test$p.value, scientific = TRUE), "\n\n")
```

P-value = 2.440503e-05

The Distribution of P-Values (Devore Example 9.22)

P-Values Are Random Variables

The P-value is computed from the test statistic, which is computed from sample data. Different samples yield different test statistics and therefore **different P-values**.

P-Values Are Random Variables

The P-value is computed from the test statistic, which is computed from sample data. Different samples yield different test statistics and therefore **different P-values**.

Key theoretical results:

- When H_0 is **true**: the P-value has a **Uniform(0, 1)** distribution
- When H_0 is **false**: the P-value distribution is **concentrated near 0**

P-Values Are Random Variables

The P-value is computed from the test statistic, which is computed from sample data. Different samples yield different test statistics and therefore **different P-values**.

Key theoretical results:

- When H_0 is **true**: the P-value has a **Uniform(0, 1)** distribution
- When H_0 is **false**: the P-value distribution is **concentrated near 0**

! Why Uniform Under H_0 ?

If H_0 is true, then $P(\text{P-value} \leq \alpha) = \alpha$ for any $\alpha \in [0, 1]$. This is precisely the definition of a Uniform(0, 1) distribution!

Recap: P-Value Uniformity via the Probability Integral Transform

Recall our earlier result: if T has continuous CDF F_T under H_0 , then $P = 1 - F_T(T) \sim \text{Uniform}(0, 1)$.

Recap: P-Value Uniformity via the Probability Integral Transform

Recall our earlier result: if T has continuous CDF F_T under H_0 , then $P = 1 - F_T(T) \sim \text{Uniform}(0, 1)$.

For the z test specifically:

Let $Z = \frac{\bar{X} - \mu_0}{\sigma/\sqrt{n}} \sim N(0, 1)$ under H_0 . The P-value for an upper-tailed test is:

$$P = 1 - \Phi(Z)$$

Recap: P-Value Uniformity via the Probability Integral Transform

Recall our earlier result: if T has continuous CDF F_T under H_0 , then $P = 1 - F_T(T) \sim \text{Uniform}(0, 1)$.

For the z test specifically:

Let $Z = \frac{\bar{X} - \mu_0}{\sigma/\sqrt{n}} \sim N(0, 1)$ under H_0 . The P-value for an upper-tailed test is:

$$P = 1 - \Phi(Z)$$

Since Φ is the CDF of Z , this is exactly the survival function form, and the probability integral transformation gives us $P \sim \text{Uniform}(0, 1)$.

What Happens Under the Alternative?

Under $H_a : \mu > \mu_0$, the test statistic $Z = \frac{\bar{X} - \mu_0}{\sigma/\sqrt{n}}$ is **not** standard normal.

What Happens Under the Alternative?

Under $H_a : \mu > \mu_0$, the test statistic $Z = \frac{\bar{X} - \mu_0}{\sigma/\sqrt{n}}$ is **not** standard normal.

Instead, Z has mean $\frac{\mu - \mu_0}{\sigma/\sqrt{n}} > 0$, so Z tends to be positive and large.

What Happens Under the Alternative?

Under $H_a : \mu > \mu_0$, the test statistic $Z = \frac{\bar{X} - \mu_0}{\sigma/\sqrt{n}}$ is **not** standard normal.

Instead, Z has mean $\frac{\mu - \mu_0}{\sigma/\sqrt{n}} > 0$, so Z tends to be positive and large.

Consequence for the P-value:

$$P = 1 - \Phi(Z) = P_{H_0}(Z' \geq Z)$$

When Z tends to be large and positive, $\Phi(Z) \rightarrow 1$, so $P \rightarrow 0$.

What Happens Under the Alternative?

Under $H_a : \mu > \mu_0$, the test statistic $Z = \frac{\bar{X} - \mu_0}{\sigma/\sqrt{n}}$ is **not** standard normal.

Instead, Z has mean $\frac{\mu - \mu_0}{\sigma/\sqrt{n}} > 0$, so Z tends to be positive and large.

Consequence for the P-value:

$$P = 1 - \Phi(Z) = P_{H_0}(Z' \geq Z)$$

When Z tends to be large and positive, $\Phi(Z) \rightarrow 1$, so $P \rightarrow 0$.

The farther μ is from μ_0 :

- The more Z is shifted to the right
- The larger $\Phi(Z)$ becomes
- The smaller the P-value $P = 1 - \Phi(Z)$
- The more P-values concentrate near 0

The Distribution of P-Values: A Unified View

Let $\delta = \frac{\mu - \mu_0}{\sigma / \sqrt{n}}$ be the **noncentrality parameter** (standardized effect size).

The Distribution of P-Values: A Unified View

Let $\delta = \frac{\mu - \mu_0}{\sigma/\sqrt{n}}$ be the **noncentrality parameter** (standardized effect size).

Under the true parameter value:

- $\delta = 0$ (i.e., H_0 true): $P \sim \text{Uniform}(0, 1)$
- $\delta > 0$ (upper-tailed H_a true): P has CDF $F_P(p) = 1 - \Phi(\Phi^{-1}(1 - p) - \delta)$


The Distribution of P-Values: A Unified View

Let $\delta = \frac{\mu - \mu_0}{\sigma/\sqrt{n}}$ be the **noncentrality parameter** (standardized effect size).

Under the true parameter value:

- $\delta = 0$ (i.e., H_0 true): $P \sim \text{Uniform}(0, 1)$
- $\delta > 0$ (upper-tailed H_a true): P has CDF $F_P(p) = 1 - \Phi(\Phi^{-1}(1 - p) - \delta)$

► Code



Power in Terms of P-Value Distribution

The **power** at a given alternative μ is the probability of rejecting H_0 :

$$\text{Power}(\mu) = P_{\mu}(\text{Reject } H_0) = P_{\mu}(P \leq \alpha) = F_P(\alpha)$$

Power in Terms of P-Value Distribution

The **power** at a given alternative μ is the probability of rejecting H_0 :

$$\text{Power}(\mu) = P_{\mu}(\text{Reject } H_0) = P_{\mu}(P \leq \alpha) = F_P(\alpha)$$

Using our formula for the P-value CDF under the alternative:

$$\text{Power}(\mu) = 1 - \Phi(\Phi^{-1}(1 - \alpha) - \delta) = \Phi(\delta - z_{1-\alpha})$$

where $\delta = \frac{\mu - \mu_0}{\sigma/\sqrt{n}}$ and $z_{1-\alpha} = \Phi^{-1}(1 - \alpha)$.

Power in Terms of P-Value Distribution

The **power** at a given alternative μ is the probability of rejecting H_0 :

$$\text{Power}(\mu) = P_{\mu}(\text{Reject } H_0) = P_{\mu}(P \leq \alpha) = F_P(\alpha)$$

Using our formula for the P-value CDF under the alternative:

$$\text{Power}(\mu) = 1 - \Phi(\Phi^{-1}(1 - \alpha) - \delta) = \Phi(\delta - z_{1-\alpha})$$

where $\delta = \frac{\mu - \mu_0}{\sigma/\sqrt{n}}$ and $z_{1-\alpha} = \Phi^{-1}(1 - \alpha)$.


This connects P-value theory to power calculations!

Simulation: P-Value Distribution (Devore Example 9.22)

► Code

Distribution of P-Values: 10,000 Simulated t Tests ($n = 4$)

Testing $H_0: \mu = 30$ vs $H_a: \mu > 30$ with $\sigma = 2$



What the Simulation Shows

```
1 # Proportion of P-values <= 0.05 in each scenario
2 cat("Proportion rejecting H0 at alpha = 0.05:\n")
```

Proportion rejecting H0 at alpha = 0.05:

```
1 cat("  H0 true (mu = 30): ", round(mean(pvals_H0_true <= 0.05), 3),
2     " (should be ~0.05)\n")
```

H0 true (mu = 30): 0.051 (should be ~0.05)

```
1 cat("  H0 false (mu = 31):", round(mean(pvals_mu31 <= 0.05), 3),
2     " (power at mu = 31)\n")
```

H0 false (mu = 31): 0.201 (power at mu = 31)

```
1 cat("  H0 false (mu = 32):", round(mean(pvals_mu32 <= 0.05), 3),
2     " (power at mu = 32)\n")
```

H0 false (mu = 32): 0.462 (power at mu = 32)

What the Simulation Shows

```
1 # Proportion of P-values <= 0.05 in each scenario
2 cat("Proportion rejecting H0 at alpha = 0.05:\n")
```

Proportion rejecting H0 at alpha = 0.05:

```
1 cat("  H0 true (mu = 30): ", round(mean(pvals_H0_true <= 0.05), 3),
2     " (should be ~0.05)\n")
```

H0 true (mu = 30): 0.051 (should be ~0.05)

```
1 cat("  H0 false (mu = 31):", round(mean(pvals_mu31 <= 0.05), 3),
2     " (power at mu = 31)\n")
```

H0 false (mu = 31): 0.201 (power at mu = 31)

```
1 cat("  H0 false (mu = 32):", round(mean(pvals_mu32 <= 0.05), 3),
2     " (power at mu = 32)\n")
```

H0 false (mu = 32): 0.462 (power at mu = 32)

Key insights (Devore Example 9.22):

- When H_0 is true, the P-value histogram is approximately **flat** (Uniform)
- Only ~5% of P-values fall below 0.05, confirming α controls the Type I error rate
- As μ moves farther from $\mu_0 = 30$, P-values concentrate closer to 0
- Even at $\mu = 32$, power is less than 50% because $n = 4$ is very small!

Common Misinterpretations of P-Values

What the P-Value is NOT

1. **Not** the probability that H_0 is true: $\text{P-value} \neq P(H_0 \text{ is true} \mid \text{data})$
2. **Not** the probability of making an error: it is $P(\text{data this extreme} \mid H_0 \text{ true})$
3. **Not** a measure of effect size: a small P-value can arise from a tiny effect with large n
4. A P-value of 0.05 does **not** mean there is a 5% chance H_0 is true
5. Failing to reject H_0 does **not** mean H_0 is true (absence of evidence $\neq$ evidence of absence)

Common Misinterpretations of P-Values

What the P-Value is NOT

1. **Not** the probability that H_0 is true: $\text{P-value} \neq P(H_0 \text{ is true} \mid \text{data})$
2. **Not** the probability of making an error: it is $P(\text{data this extreme} \mid H_0 \text{ true})$
3. **Not** a measure of effect size: a small P-value can arise from a tiny effect with large n
4. A P-value of 0.05 does **not** mean there is a 5% chance H_0 is true
5. Failing to reject H_0 does **not** mean H_0 is true (absence of evidence $\neq$ evidence of absence)

ASA Statement on P-Values (2016)

The American Statistical Association emphasized: “a P-value does not provide a good measure of evidence regarding a model or hypothesis.” P-values should be supplemented with effect sizes, confidence intervals, and clinical judgment.

Statistical vs Clinical Significance

► Code

```
n = 5000
```

► Code

```
Sample mean: 120.37 mmHg
```

► Code

```
P-value: 0.083
```

► Code

```
95% CI: ( 119.95 , 120.8 )
```

► Code

```
Statistically significant (P < 0.05)? FALSE
```

► Code

```
Clinically meaningful difference? A 0.3 mmHg difference
```

► Code

```
in systolic BP has NO clinical relevance!
```



Statistical vs Clinical Significance

► Code

```
n = 5000
```

► Code

```
Sample mean: 120.37 mmHg
```

► Code

```
P-value: 0.083
```

► Code

```
95% CI: ( 119.95 , 120.8 )
```

► Code

```
Statistically significant (P < 0.05)? FALSE
```

► Code

```
Clinically meaningful difference? A 0.3 mmHg difference
```

► Code

```
in systolic BP has NO clinical relevance!
```



Power for the t Test: The Noncentral t Distribution

Review: Power for the z Test (Lesson 13)

For the one-sample z test (σ known), we derived closed-form expressions:

Alternative	$\beta(\mu')$
$H_a : \mu > \mu_0$	$\Phi\left(z_\alpha + \frac{\mu_0 - \mu'}{\sigma/\sqrt{n}}\right)$
$H_a : \mu < \mu_0$	$1 - \Phi\left(-z_\alpha + \frac{\mu_0 - \mu'}{\sigma/\sqrt{n}}\right)$
$H_a : \mu \neq \mu_0$	$\Phi\left(z_{\alpha/2} + \frac{\mu_0 - \mu'}{\sigma/\sqrt{n}}\right) - \Phi\left(-z_{\alpha/2} + \frac{\mu_0 - \mu'}{\sigma/\sqrt{n}}\right)$

Review: Power for the z Test (Lesson 13)

For the one-sample z test (σ known), we derived closed-form expressions:

Alternative	$\beta(\mu')$
$H_a : \mu > \mu_0$	$\Phi\left(z_\alpha + \frac{\mu_0 - \mu'}{\sigma/\sqrt{n}}\right)$
$H_a : \mu < \mu_0$	$1 - \Phi\left(-z_\alpha + \frac{\mu_0 - \mu'}{\sigma/\sqrt{n}}\right)$
$H_a : \mu \neq \mu_0$	$\Phi\left(z_{\alpha/2} + \frac{\mu_0 - \mu'}{\sigma/\sqrt{n}}\right) - \Phi\left(-z_{\alpha/2} + \frac{\mu_0 - \mu'}{\sigma/\sqrt{n}}\right)$

Limitation: These formulas require σ to be known. In practice, we use the t test, and the analysis of power becomes more complex.

The Noncentral t Distribution (Devore 9.2)

When H_0 is **false** (i.e., $\mu = \mu' \neq \mu_0$), the t test statistic

$$T = \frac{\bar{X} - \mu_0}{S/\sqrt{n}}$$

does **not** follow a t_{n-1} distribution. Instead, it follows a **noncentral t distribution**.

The Noncentral t Distribution (Devore 9.2)

When H_0 is false (i.e., $\mu = \mu' \neq \mu_0$), the t test statistic

$$T = \frac{\bar{X} - \mu_0}{S/\sqrt{n}}$$

does **not** follow a t_{n-1} distribution. Instead, it follows a **noncentral t distribution**.

! Definition (Devore 9.2)

Let $Z \sim N(0, 1)$ and $Y \sim \chi_m^2$ be independent. For any real number δ , the random variable

$$\frac{Z + \delta}{\sqrt{Y/m}}$$

has a **noncentral t distribution** with m degrees of freedom and **noncentrality parameter** δ .

When $\delta = 0$, this reduces to the ordinary (central) t_m distribution.

Noncentrality Parameter for the One-Sample t Test

When $\mu = \mu'$ (the true mean under H_a), the one-sample t statistic has a noncentral t distribution with:

- Degrees of freedom: $n - 1$
- Noncentrality parameter: $\delta = \frac{\mu' - \mu_0}{\sigma / \sqrt{n}}$

Noncentrality Parameter for the One-Sample t Test

When $\mu = \mu'$ (the true mean under H_a), the one-sample t statistic has a noncentral t distribution with:

- **Degrees of freedom:** $n - 1$
- **Noncentrality parameter:** $\delta = \frac{\mu' - \mu_0}{\sigma / \sqrt{n}}$

The noncentrality parameter δ measures how far the true mean is from the null value, in units of the standard error.

Noncentrality Parameter for the One-Sample t Test

When $\mu = \mu'$ (the true mean under H_a), the one-sample t statistic has a noncentral t distribution with:

- Degrees of freedom: $n - 1$
- Noncentrality parameter: $\delta = \frac{\mu' - \mu_0}{\sigma / \sqrt{n}}$

The noncentrality parameter δ measures how far the true mean is from the null value, in units of the standard error.

```
1 # Example: mu_0 = 120, mu' = 125, sigma = 15, n = 25
2 mu0 <- 120; mu_prime <- 125; sigma <- 15; n <- 25
3 delta <- (mu_prime - mu0) / (sigma / sqrt(n))
4 cat("Noncentrality parameter: delta =", round(delta, 3))
```


Noncentrality parameter: delta = 1.667

Visualizing the Noncentral t Distribution

► Code

Noncentral t Distribution (df = 9)

Larger δ shifts the distribution right → more power for upper-tailed tests



Computing Power with the Noncentral t

Power for an upper-tailed test $H_a : \mu > \mu_0$:

$$\text{Power} = P(T \geq t_{\alpha, n-1}) = 1 - F(t_{\alpha, n-1}; n - 1, \delta)$$

where $F(x; m, \delta)$ is the CDF of the noncentral t distribution.

Computing Power with the Noncentral t

Power for an upper-tailed test $H_a : \mu > \mu_0$:

$$\text{Power} = P(T \geq t_{\alpha, n-1}) = 1 - F(t_{\alpha, n-1}; n - 1, \delta)$$

where $F(x; m, \delta)$ is the CDF of the noncentral t distribution.

In R: `pt(t_crit, df, ncp = delta, lower.tail = FALSE)`

Computing Power with the Noncentral t

Power for an upper-tailed test $H_a : \mu > \mu_0$:

$$\text{Power} = P(T \geq t_{\alpha, n-1}) = 1 - F(t_{\alpha, n-1}; n - 1, \delta)$$

where $F(x; m, \delta)$ is the CDF of the noncentral t distribution.

In R: `pt(t_crit, df, ncp = delta, lower.tail = FALSE)`

```
1 # Example (Devore Example 9.14 adapted to medical context):
2 # Testing whether mean systolic BP exceeds 120 mmHg
3 # H0: mu = 120, Ha: mu > 120, alpha = 0.05, n = 10, sigma = 10
4 # What is the power when mu' = 125?
5 n <- 10; mu0 <- 120; mu_prime <- 125; sigma <- 10; alpha <- 0.05
6
7 delta <- (mu_prime - mu0) / (sigma / sqrt(n))
8 t_crit <- qt(1 - alpha, df = n - 1)
9
10 power <- pt(t_crit, df = n - 1, ncp = delta, lower.tail = FALSE)
11 cat("Noncentrality parameter: delta =", round(delta, 3), "\n")
```


Noncentrality parameter: delta = 1.581

```
1 cat("Critical value: t_{0.05, 9} =", round(t_crit, 3), "\n")
```

Critical value: $t_{\{0.05, 9\}} = 1.833$

Power Curves for the t Test: Medical Example

► Code



Key Features of Power Curves

Two intuitive features:

1. For any **fixed** departure from H_0 , power **increases** with sample size
2. For any **fixed** sample size, power **increases** as μ' moves farther from μ_0

Key Features of Power Curves

Two intuitive features:

1. For any **fixed** departure from H_0 , power **increases** with sample size
2. For any **fixed** sample size, power **increases** as μ' moves farther from μ_0

Factors that increase power:

- **Larger n :** more data $\Rightarrow$ smaller SE $\Rightarrow$ easier to detect departures
- **Larger effect size $|\mu' - \mu_0|$:** bigger departures are easier to detect
- **Smaller σ :** less noise $\Rightarrow$ easier to detect signal
- **Larger α :** easier to reject (but at the cost of more Type I errors)

Key Features of Power Curves

Two intuitive features:

1. For any **fixed** departure from H_0 , power **increases** with sample size
2. For any **fixed** sample size, power **increases** as μ' moves farther from μ_0

Factors that increase power:

- **Larger n** : more data $\Rightarrow$ smaller SE $\Rightarrow$ easier to detect departures
- **Larger effect size $|\mu' - \mu_0|$** : bigger departures are easier to detect
- **Smaller σ** : less noise $\Rightarrow$ easier to detect signal
- **Larger α** : easier to reject (but at the cost of more Type I errors)

Clinical Trial Implication

The power analysis must be conducted **before** collecting data. It ensures the study has sufficient sample size to detect a **clinically meaningful** effect.

Sample Size Determination

Why Determine Sample Size?

Clinical scenario: You are designing a randomized trial to test whether a new antihypertensive drug lowers systolic blood pressure more than a placebo.

Why Determine Sample Size?

Clinical scenario: You are designing a randomized trial to test whether a new antihypertensive drug lowers systolic blood pressure more than a placebo.

Key questions:

- How many patients do we need to **detect** a meaningful effect?
- What effect size is **clinically important**?
- What balance of Type I and Type II errors is acceptable?

Why Determine Sample Size?

Clinical scenario: You are designing a randomized trial to test whether a new antihypertensive drug lowers systolic blood pressure more than a placebo.

Key questions:

- How many patients do we need to **detect** a meaningful effect?
- What effect size is **clinically important**?
- What balance of Type I and Type II errors is acceptable?

! The Three Ingredients for Sample Size

1. **Significance level** α (usually 0.05)
2. **Desired power** $1 - \beta$ (usually 0.80 or 0.90)
3. **Clinically meaningful effect size** $\Delta = |\mu' - \mu_0|$ (determined by clinical judgment)
4. **(Estimated) standard deviation** σ

Sample Size Formula for the z Test (Review)

For the one-sample z test, the required sample size for power $1 - \beta$ at alternative μ' :

$$n = \begin{cases} \left[\frac{\sigma(z_\alpha + z_\beta)}{\mu_0 - \mu'} \right]^2 & \text{one-tailed test} \\ \left[\frac{\sigma(z_{\alpha/2} + z_\beta)}{\mu_0 - \mu'} \right]^2 & \text{two-tailed test (approximate)} \end{cases}$$

Sample Size Formula for the z Test (Review)

For the one-sample z test, the required sample size for power $1 - \beta$ at alternative μ' :

$$n = \begin{cases} \left[\frac{\sigma(z_\alpha + z_\beta)}{\mu_0 - \mu'} \right]^2 & \text{one-tailed test} \\ \left[\frac{\sigma(z_{\alpha/2} + z_\beta)}{\mu_0 - \mu'} \right]^2 & \text{two-tailed test (approximate)} \end{cases}$$

```
1 # How many patients to detect a 5 mmHg reduction with 80% power?
2 delta <- 5 # clinically meaningful difference (mmHg)
3 sigma <- 15 # estimated SD
4 alpha <- 0.05
5 beta <- 0.20 # power = 0.80
6
7 # Two-tailed z test
8 n_z <- (sigma * (qnorm(1 - alpha/2) + qnorm(1 - beta)) / delta)^2
9 cat("z test sample size (two-tailed):", ceiling(n_z), "patients")
```

z test sample size (two-tailed): 71 patients

R's `power.t.test()` Function

R provides `power.t.test()` for sample size and power calculations for the one-sample and two-sample t tests. It accounts for the noncentral t distribution automatically.

R's `power.t.test()` Function

R provides `power.t.test()` for sample size and power calculations for the one-sample and two-sample t tests. It accounts for the noncentral t distribution automatically.

```
1 # Same example using power.t.test()
2 result <- power.t.test(
3   delta = 5,           # difference to detect
4   sd = 15,             # estimated standard deviation
5   sig.level = 0.05,    # alpha
6   power = 0.80,        # desired power
7   type = "one.sample", # one-sample test
8   alternative = "two.sided"
9 )
10 result
```

One-sample t test power calculation

```
      n = 72.58407
  delta = 5
    sd = 15
sig.level = 0.05
  power = 0.8
alternative = two.sided
```

R's `power.t.test()` Function

R provides `power.t.test()` for sample size and power calculations for the one-sample and two-sample t tests. It accounts for the noncentral t distribution automatically.

```
1 # Same example using power.t.test()
2 result <- power.t.test(
3   delta = 5,           # difference to detect
4   sd = 15,             # estimated standard deviation
5   sig.level = 0.05,    # alpha
6   power = 0.80,        # desired power
7   type = "one.sample", # one-sample test
8   alternative = "two.sided"
9 )
10 result
```

One-sample t test power calculation

```
      n = 72.58407
delta = 5
sd = 15
sig.level = 0.05
power = 0.8
alternative = two.sided
```

```
1 cat("Required sample size:", ceiling(result$n), "patients")
```

Example: Planning a Clinical Trial for Blood Pressure

Setting: Researchers are planning a trial to test whether a new drug reduces systolic blood pressure. Historical data suggests $\sigma \approx 12$ mmHg. A reduction of 5 mmHg is considered clinically meaningful.

```
1 # One-sided test: H0: mu >= 0 vs Ha: mu < 0 (reduction in BP)
2 power_result <- power.t.test(
3   delta = 5,
4   sd = 12,
5   sig.level = 0.05,
6   power = 0.80,
7   type = "one.sample",
8   alternative = "one.sided"
9 )
10 cat("For 80% power:", ceiling(power_result$n), "patients needed\n")
```

For 80% power: 38 patients needed

```
1 # What about 90% power?
2 power_90 <- power.t.test(
3   delta = 5, sd = 12, sig.level = 0.05,
4   power = 0.90, type = "one.sample", alternative = "one.sided"
5 )
6 cat("For 90% power:", ceiling(power_90$n), "patients needed\n")
```

For 90% power: 51 patients needed


```
1 # What about a smaller effect (3 mmHg)?
```


Sample Size as a Function of Effect Size and Power

► Code


Sample Size vs Effect Size for One-Sample t Test

$\sigma = 12$ mmHg, $\alpha = 0.05$, two-tailed



Power vs Sample Size

$\sigma = 12$ mmHg, $\alpha = 0.05$, two-tailed



Example: Neonatal ICU Glucose Monitoring

Setting: A neonatal ICU has historically maintained mean blood glucose at 80 mg/dL (sd = 20 mg/dL) in premature infants. A new feeding protocol is introduced. The clinical team wants to detect any shift in mean glucose of at least 8 mg/dL with 90% power.

```
1 # How many infants are needed?
2 nicu_power <- power.t.test(
3   delta = 8,
4   sd = 20,
5   sig.level = 0.05,
6   power = 0.90,
7   type = "one.sample",
8   alternative = "two.sided"
9 )
10
11 cat("=== NICU Study Design ===\n")
```

=== NICU Study Design ===

```
1 cat("Effect to detect: 8 mg/dL\n")
```

Effect to detect: 8 mg/dL

```
1 cat("Assumed SD: 20 mg/dL\n")
```

Assumed SD: 20 mg/dL

```
1 cat("Significance level: 0.05\n")
```

Manual Power Calculation with the Noncentral t

Let's verify `power.t.test()` using the noncentral t distribution directly:

```
1 # NICU example: n = 70, delta = 8, sigma = 20, alpha = 0.05, two-tailed
2 n <- ceiling(nicu_power$n)
3 mu0 <- 80; mu_prime <- 88; sigma <- 20
4 alpha <- 0.05
5
6 # Noncentrality parameter
7 delta_ncp <- (mu_prime - mu0) / (sigma / sqrt(n))
8 cat("Noncentrality parameter: delta =", round(delta_ncp, 3), "\n")
```

Noncentrality parameter: delta = 3.298

```
1 # Critical values for two-tailed test
2 t_crit <- qt(1 - alpha/2, df = n - 1)
3 cat("Critical value: t_{0.025, ", n-1, "} =", round(t_crit, 3), "\n")
```

Critical value: $t_{\{0.025, 67\}} = 1.996$

```
1 # Power = P(|T| >= t_crit) under noncentral t
2 power_upper <- pt(t_crit, df = n - 1, ncp = delta_ncp, lower.tail = FALSE)
3 power_lower <- pt(-t_crit, df = n - 1, ncp = delta_ncp, lower.tail = TRUE)
4 power_total <- power_upper + power_lower
5
6 cat("Power (upper tail):", round(power_upper, 4), "\n")
```

Power (upper tail): 0.9016





Connecting P-Values, Power, and Sample Size

The Complete Picture

► Code

P-value Distributions When H_0 Is False ($\mu = 125$)

$H_0: \mu = 120$, two-tailed, $\sigma = 15$



Summary: P-Values and Their Relationship to Testing

Concept	What It Tells You
P-value	Strength of evidence against H_0 (smaller = stronger)
α	Threshold for “significant” evidence (set in advance)
$\beta(\mu')$	Probability of missing a true effect of size μ'
Power = $1 - \beta$	Probability of detecting a true effect
Sample size n	Number of observations needed for desired power

Summary: P-Values and Their Relationship to Testing

Concept	What It Tells You
P-value	Strength of evidence against H_0 (smaller = stronger)
α	Threshold for “significant” evidence (set in advance)
$\beta(\mu')$	Probability of missing a true effect of size μ'
Power = $1 - \beta$	Probability of detecting a true effect
Sample size n	Number of observations needed for desired power

The chain of reasoning:

Larger $n \implies$ Smaller SE $\implies$ Larger $|T| \implies$ Smaller P-values $\implies$ More Power

Practical Guidelines for Medical Research

Best Practices

1. **Always report P-values AND confidence intervals** — the CI tells you about effect size
2. **Conduct power analysis BEFORE the study** — not after!
3. **Use clinically meaningful effect sizes** — not just “any” difference
4. **Don’t p-hack:** pre-register hypotheses and analysis plans
5. **Interpret in context:** a P-value of 0.048 vs 0.052 should not lead to radically different conclusions
6. **Consider multiple comparisons:** if testing many hypotheses, adjust α (Bonferroni, FDR, etc.)

Practical Guidelines for Medical Research

Best Practices

1. **Always report P-values AND confidence intervals** – the CI tells you about effect size
2. **Conduct power analysis BEFORE the study** – not after!
3. **Use clinically meaningful effect sizes** – not just “any” difference
4. **Don’t p-hack:** pre-register hypotheses and analysis plans
5. **Interpret in context:** a P-value of 0.048 vs 0.052 should not lead to radically different conclusions
6. **Consider multiple comparisons:** if testing many hypotheses, adjust α (Bonferroni, FDR, etc.)

FDA Guidance

Clinical trials typically require $\alpha = 0.05$ (two-sided) and power ≥ 0.80 . Phase III trials often require power ≥ 0.90 to ensure adequate evidence for drug approval.

Multiple Testing and P-Value Adjustment

The Multiple Testing Problem

Scenario: A genomics study tests 20,000 genes for differential expression between tumor and normal tissue. Using $\alpha = 0.05$ for each test...

The Multiple Testing Problem

Scenario: A genomics study tests 20,000 genes for differential expression between tumor and normal tissue. Using $\alpha = 0.05$ for each test...

Expected false positives under H_0 (no true effects):

$$\text{Expected false discoveries} = m \times \alpha = 20,000 \times 0.05 = 1,000$$

The Multiple Testing Problem

Scenario: A genomics study tests 20,000 genes for differential expression between tumor and normal tissue. Using $\alpha = 0.05$ for each test...

Expected false positives under H_0 (no true effects):

$$\text{Expected false discoveries} = m \times \alpha = 20,000 \times 0.05 = 1,000$$

Even if no genes are truly differentially expressed, we expect ~1,000 “significant” results by chance alone!

The Multiple Testing Problem

Scenario: A genomics study tests 20,000 genes for differential expression between tumor and normal tissue. Using $\alpha = 0.05$ for each test...

Expected false positives under H_0 (no true effects):

$$\text{Expected false discoveries} = m \times \alpha = 20,000 \times 0.05 = 1,000$$

Even if no genes are truly differentially expressed, we expect ~1,000 “significant” results by chance alone!

The Core Problem

When conducting m hypothesis tests, the probability of making **at least one** Type I error grows rapidly with m , even if each individual test has size α .

Family-Wise Error Rate (FWER)

Definition

The **family-wise error rate (FWER)** is the probability of making one or more Type I errors among all m tests:

$$\text{FWER} = P(\text{at least one false rejection})$$

Family-Wise Error Rate (FWER)

Definition

The **family-wise error rate (FWER)** is the probability of making one or more Type I errors among all m tests:

$$\text{FWER} = P(\text{at least one false rejection})$$

If the m tests are **independent** and all null hypotheses are true:

$$\text{FWER} = 1 - P(\text{no false rejections}) = 1 - (1 - \alpha)^m$$

Family-Wise Error Rate (FWER)

Definition

The **family-wise error rate (FWER)** is the probability of making one or more Type I errors among all m tests:

$$\text{FWER} = P(\text{at least one false rejection})$$

If the m tests are **independent** and all null hypotheses are true:

$$\text{FWER} = 1 - P(\text{no false rejections}) = 1 - (1 - \alpha)^m$$

```
1 # FWER as a function of m for alpha = 0.05
2 m_tests <- c(1, 2, 5, 10, 20, 50, 100)
3 fwer <- 1 - (1 - 0.05)^m_tests
4 tibble(m = m_tests, FWER = round(fwer, 3)) |>
5   pivot_wider(names_from = m, values_from = FWER) |>
6   knitr::kable(col.names = paste0("m = ", m_tests))
```


m = 1	m = 2	m = 5	m = 10	m = 20	m = 50	m = 100
0.05	0.098	0.226	0.401	0.642	0.923	0.994

Visualizing the FWER Inflation

► Code

FWER Inflation with Multiple Testing

Using individual $\alpha = 0.05$ for each test (assuming independence)



The Bonferroni Correction

! Bonferroni Correction

To control the FWER at level α when conducting m tests, reject each individual hypothesis only if its P-value satisfies:

$$p_i \leq \frac{\alpha}{m}$$

Equivalently, compute **adjusted P-values**: $\tilde{p}_i = \min(m \cdot p_i, 1)$

The Bonferroni Correction

! Bonferroni Correction

To control the FWER at level α when conducting m tests, reject each individual hypothesis only if its P-value satisfies:

$$p_i \leq \frac{\alpha}{m}$$

Equivalently, compute **adjusted P-values**: $\tilde{p}_i = \min(m \cdot p_i, 1)$

Theoretical justification (Boole's inequality):

$$\text{FWER} = P\left(\bigcup_{i \in \mathcal{H}_0} \{\text{Reject } H_{0i}\}\right) \leq \sum_{i \in \mathcal{H}_0} P(\text{Reject } H_{0i}) \leq m_0 \cdot \frac{\alpha}{m} \leq \alpha$$

where m_0 is the number of true null hypotheses.

Bonferroni: Example

Clinical trial: Testing a drug's effect on 5 different biomarkers.

```
1 # Raw P-values from 5 tests
2 biomarkers <- c("CRP", "IL-6", "TNF-alpha", "Cortisol", "Fibrinogen")
3 raw_pvalues <- c(0.008, 0.023, 0.041, 0.067, 0.12)
4
5 # Bonferroni adjustment
6 m <- length(raw_pvalues)
7 bonf_adjusted <- pmin(raw_pvalues * m, 1) # p.adjust does this
8
9 # Using R's built-in function
10 bonf_adjusted_r <- p.adjust(raw_pvalues, method = "bonferroni")
11
12 tibble(
13   Biomarker = biomarkers,
14   `Raw P-value` = raw_pvalues,
15   `Bonferroni Adjusted` = round(bonf_adjusted_r, 3),
16   `Significant (alpha=0.05)` = ifelse(bonf_adjusted_r <= 0.05, "Yes", "No")
17 ) |> knitr::kable()
```

Biomarker	Raw P-value	Bonferroni Adjusted	Significant (alpha=0.05)
CRP	0.008	0.040	Yes
IL-6	0.023	0.115	No
TNF-alpha	0.041	0.205	No
Cortisol	0.067	0.335	No
Fibrinogen	0.12	0.6	No

Bonferroni: Example

Clinical trial: Testing a drug's effect on 5 different biomarkers.

```
1 # Raw P-values from 5 tests
2 biomarkers <- c("CRP", "IL-6", "TNF-alpha", "Cortisol", "Fibrinogen")
3 raw_pvalues <- c(0.008, 0.023, 0.041, 0.067, 0.12)
4
5 # Bonferroni adjustment
6 m <- length(raw_pvalues)
7 bonf_adjusted <- pmin(raw_pvalues * m, 1) # p.adjust does this
8
9 # Using R's built-in function
10 bonf_adjusted_r <- p.adjust(raw_pvalues, method = "bonferroni")
11
12 tibble(
13   Biomarker = biomarkers,
14   `Raw P-value` = raw_pvalues,
15   `Bonferroni Adjusted` = round(bonf_adjusted_r, 3),
16   `Significant (alpha=0.05)` = ifelse(bonf_adjusted_r <= 0.05, "Yes", "No")
17 ) |> knitr::kable()
```

Biomarker	Raw P-value	Bonferroni Adjusted	Significant (alpha=0.05)
CRP	0.008	0.040	Yes
IL-6	0.023	0.115	No
TNF-alpha	0.041	0.205	No
Cortisol	0.067	0.335	No
Fibrinogen	0.12	0.6	No

Limitations of Bonferroni

Bonferroni is Conservative

The Bonferroni correction controls FWER at level **at most** α , but often the actual FWER is much smaller, leading to reduced **power**.

Limitations of Bonferroni

Bonferroni is Conservative

The Bonferroni correction controls FWER at level **at most** α , but often the actual FWER is much smaller, leading to reduced **power**.

When Bonferroni is too conservative:

- Tests are positively correlated (common in genomics)
- Large number of tests ($m > 100$)
- Many true alternative hypotheses

Limitations of Bonferroni

Bonferroni is Conservative

The Bonferroni correction controls FWER at level **at most** α , but often the actual FWER is much smaller, leading to reduced **power**.

When Bonferroni is too conservative:

- Tests are positively correlated (common in genomics)
- Large number of tests ($m > 100$)
- Many true alternative hypotheses

The problem: With $m = 20,000$ tests, Bonferroni requires $p < 0.05/20,000 = 2.5 \times 10^{-6}$. This may miss many true effects.

Alternative: False Discovery Rate (FDR)

Instead of controlling the probability of **any** false positives (FWER), we can control the **proportion** of false positives among all discoveries.

Definition: False Discovery Rate

Let V = number of false discoveries (Type I errors) and R = total number of rejections. The **false discovery rate** is:

$$\text{FDR} = E \left[\frac{V}{R} \mid R > 0 \right] \cdot P(R > 0)$$

Often approximated as $E[V/R]$ with the convention that $V/R = 0$ when $R = 0$.

Alternative: False Discovery Rate (FDR)

Instead of controlling the probability of **any** false positives (FWER), we can control the **proportion** of false positives among all discoveries.

Definition: False Discovery Rate

Let V = number of false discoveries (Type I errors) and R = total number of rejections. The **false discovery rate** is:

$$\text{FDR} = E \left[\frac{V}{R} \mid R > 0 \right] \cdot P(R > 0)$$

Often approximated as $E[V/R]$ with the convention that $V/R = 0$ when $R = 0$.

Interpretation: FDR = 0.05 means that, on average, 5% of our “discoveries” are false positives. This is often acceptable in exploratory research.

Benjamini-Hochberg (BH) Procedure

! Benjamini-Hochberg Procedure (1995)

To control FDR at level q :

1. Order the m P-values: $p_{(1)} \leq p_{(2)} \leq \dots \leq p_{(m)}$
2. Find the largest k such that $p_{(k)} \leq \frac{k}{m} q$
3. Reject all hypotheses with $p_{(i)} \leq p_{(k)}$ (i.e., reject $H_{(1)}, \dots, H_{(k)}$)

Benjamini-Hochberg (BH) Procedure

! Benjamini-Hochberg Procedure (1995)

To control FDR at level q :

1. Order the m P-values: $p_{(1)} \leq p_{(2)} \leq \dots \leq p_{(m)}$
2. Find the largest k such that $p_{(k)} \leq \frac{k}{m} q$
3. Reject all hypotheses with $p_{(i)} \leq p_{(k)}$ (i.e., reject $H_{(1)}, \dots, H_{(k)}$)

Adjusted P-values: The BH-adjusted P-value for the i -th ordered hypothesis is:

$$\tilde{p}_{(i)} = \min_{j \geq i} \left\{ \frac{m}{j} p_{(j)} \right\}$$


(enforcing monotonicity: larger raw P-values cannot have smaller adjusted P-values)

Visualizing the BH Procedure

► Code

Benjamini-Hochberg Procedure (FDR = 0.10)

Reject all hypotheses with rank $\leq k^*$, where k^* is the largest k with $p(k) \leq (k/m)q$



BH Procedure: Example

```
1 # Same biomarker example
2 biomarkers <- c("CRP", "IL-6", "TNF-alpha", "Cortisol", "Fibrinogen")
3 raw_pvalues <- c(0.008, 0.023, 0.041, 0.067, 0.12)
4
5 # Multiple adjustment methods
6 adjustments <- tibble(
7   Biomarker = biomarkers,
8   `Raw P` = raw_pvalues,
9   Bonferroni = p.adjust(raw_pvalues, method = "bonferroni"),
10  Holm = p.adjust(raw_pvalues, method = "holm"),
11  BH = p.adjust(raw_pvalues, method = "BH"),
12  BY = p.adjust(raw_pvalues, method = "BY")
13 )
14
15 adjustments |>
16   mutate(across(where(is.numeric), ~round(., 3))) |>
17   knitr::kable()
```

Biomarker	Raw P	Bonferroni	Holm	BH	BY
CRP	0.008	0.040	0.040	0.040	0.091
IL-6	0.023	0.115	0.092	0.057	0.131
TNF-alpha	0.041	0.205	0.122	0.067	0.154
Cortisol	0.067	0.335	0.183	0.083	0.167
Fibrinogen	0.12	0.60	0.36	0.12	0.18

BH Procedure: Example

```
1 # Same biomarker example
2 biomarkers <- c("CRP", "IL-6", "TNF-alpha", "Cortisol", "Fibrinogen")
3 raw_pvalues <- c(0.008, 0.023, 0.041, 0.067, 0.12)
4
5 # Multiple adjustment methods
6 adjustments <- tibble(
7   Biomarker = biomarkers,
8   `Raw P` = raw_pvalues,
9   Bonferroni = p.adjust(raw_pvalues, method = "bonferroni"),
10  Holm = p.adjust(raw_pvalues, method = "holm"),
11  BH = p.adjust(raw_pvalues, method = "BH"),
12  BY = p.adjust(raw_pvalues, method = "BY")
13 )
14
15 adjustments |>
16   mutate(across(where(is.numeric), ~round(., 3))) |>
17   knitr::kable()
```

Biomarker	Raw P	Bonferroni	Holm	BH	BY
CRP	0.008	0.040	0.040	0.040	0.091
IL-6	0.023	0.115	0.092	0.057	0.131
TNF-alpha	0.041	0.205	0.122	0.067	0.154
Cortisol	0.067	0.335	0.183	0.083	0.167
Fibrinogen	0.12	0.60	0.36	0.12	0.18

Comparison of Adjustment Methods

Method	Controls	Formula	Conservative?
Bonferroni	FWER	$\tilde{p}_i = m \cdot p_i$	Very
Holm	FWER	Step-down Bonferroni	Less than Bonferroni
Hochberg	FWER	Step-up Bonferroni	Less than Holm
BH	FDR	$\tilde{p}_{(i)} = \min_{j \geq i} \frac{m}{j} p_{(j)}$	Not very
BY	FDR (any dependence)	BH with correction factor	More than BH

Comparison of Adjustment Methods

Method	Controls	Formula	Conservative?
Bonferroni	FWER	$\tilde{p}_i = m \cdot p_i$	Very
Holm	FWER	Step-down Bonferroni	Less than Bonferroni
Hochberg	FWER	Step-up Bonferroni	Less than Holm
BH	FDR	$\tilde{p}_{(i)} = \min_{j \geq i} \frac{m}{j} p_{(j)}$	Not very
BY	FDR (any dependence)	BH with correction factor	More than BH

 Practical Guidance

- Use **Bonferroni** or **Holm** when controlling FWER is essential (confirmatory studies, regulatory submissions)
- Use **BH** for exploratory research where some false discoveries are acceptable (genomics, screening)
- Use **BY** when tests may have arbitrary dependence structure

Multiple Testing in R: `p.adjust()`


```
1 # Simulate a genomics-style analysis
2 set.seed(551)
3 m <- 1000 # 1000 genes
4 n_true_effects <- 50 # 50 genes truly differentially expressed
5
6 # Generate P-values
7 pvals <- c(
8   rbeta(n_true_effects, 0.5, 10), # True effects: small P-values
9   runif(m - n_true_effects)       # Null: uniform P-values
10 )
11
12 # Apply different correction methods
13 results <- tibble(
14   raw = pvals,
15   bonferroni = p.adjust(pvals, method = "bonferroni"),
16   holm = p.adjust(pvals, method = "holm"),
17   BH = p.adjust(pvals, method = "BH")
18 )
19
20 # Count discoveries at alpha/q = 0.05
21 cat("Number of 'significant' results (adjusted P <= 0.05):\n")
```

Number of 'significant' results (adjusted P <= 0.05):

Simulation: Comparing FWER and FDR Control

► Code

Comparison of Bonferroni (FWER control) vs BH (FDR control)
Simulation: $m = 100$ tests, 10% true alternatives, 1000 replications



When to Use Which Correction?

! Decision Framework

Use FWER control (Bonferroni/Holm) when:

- False positives have serious consequences (e.g., drug approval)
- Small number of pre-specified hypotheses
- Confirmatory analysis with regulatory implications

Use FDR control (BH) when:

- Exploratory analysis where findings will be validated
- Large number of tests (genomics, neuroimaging)
- Some false discoveries are acceptable in exchange for more power
- Generating hypotheses for future studies

When to Use Which Correction?

! Decision Framework

Use FWER control (Bonferroni/Holm) when:

- False positives have serious consequences (e.g., drug approval)
- Small number of pre-specified hypotheses
- Confirmatory analysis with regulatory implications

Use FDR control (BH) when:

- Exploratory analysis where findings will be validated
- Large number of tests (genomics, neuroimaging)
- Some false discoveries are acceptable in exchange for more power
- Generating hypotheses for future studies

Medical research examples:

- **Phase III clinical trial** (primary endpoint): Bonferroni or no adjustment (single test)
- **Genome-wide association study**: BH or even less stringent (will be replicated)
- **Multiple secondary endpoints**: Holm or Hochberg
- **Biomarker discovery panel**: BH

Multiple Testing: Summary

1. **The problem:** Multiple tests inflate the chance of false discoveries
2. **FWER** = $P(\text{at least one false positive})$; controlled by **Bonferroni, Holm**
3. **FDR** = $E(\text{proportion of false discoveries})$; controlled by **BH, BY**
4. **Trade-off:** Stricter error control $\Rightarrow$ lower power
5. **In R:** `p.adjust(pvalues, method = "bonferroni")` or `method = "BH"`

Multiple Testing: Summary

1. **The problem:** Multiple tests inflate the chance of false discoveries
2. **FWER** = $P(\text{at least one false positive})$; controlled by **Bonferroni, Holm**
3. **FDR** = $E(\text{proportion of false discoveries})$; controlled by **BH, BY**
4. **Trade-off:** Stricter error control $\Rightarrow$ lower power
5. **In R:** `p.adjust(pvalues, method = "bonferroni")` or `method = "BH"`

Key Insight

Multiple testing adjustment is not optional—it's a fundamental requirement for honest science. Always report whether and how P-values were adjusted.

Key Takeaways and Looking Ahead

Key Takeaways

! Today's Main Points

1. **P-value definition:** probability of observing data as extreme or more extreme, assuming H_0 is true
2. **P-value uniformity:** Under H_0 , P-values are Uniform(0,1) by the probability integral transformation
3. **Size of a test:** P-value tests have exact size α because $P(P \leq \alpha) = \alpha$ when $P \sim \text{Uniform}(0, 1)$
4. **Conditional P-values:** Fisher's exact test conditions on sufficient statistics to eliminate nuisance parameters
5. **Noncentral t:** when H_0 is false, T follows a noncentral t with $\delta = (\mu' - \mu_0)/(\sigma/\sqrt{n})$
6. **Power in R:** `pt(t_crit, df, ncp = delta)` for manual; `power.t.test()` for convenience
7. **Multiple testing:** FWER controlled by Bonferroni/Holm; FDR controlled by BH procedure

Key Takeaways

! Today's Main Points

1. **P-value definition:** probability of observing data as extreme or more extreme, assuming H_0 is true
2. **P-value uniformity:** Under H_0 , P-values are Uniform(0,1) by the probability integral transformation
3. **Size of a test:** P-value tests have exact size α because $P(P \leq \alpha) = \alpha$ when $P \sim \text{Uniform}(0, 1)$
4. **Conditional P-values:** Fisher's exact test conditions on sufficient statistics to eliminate nuisance parameters
5. **Noncentral t:** when H_0 is false, T follows a noncentral t with $\delta = (\mu' - \mu_0)/(\sigma/\sqrt{n})$
6. **Power in R:** `pt(t_crit, df, ncp = delta)` for manual; `power.t.test()` for convenience
7. **Multiple testing:** FWER controlled by Bonferroni/Holm; FDR controlled by BH procedure

Common mistakes to avoid:

- Interpreting P-value as $P(H_0 \text{ true})$
- Confusing statistical and clinical significance
- Computing power **after** seeing the data (post hoc power)
- Failing to adjust for multiple comparisons in exploratory analyses
- Using Bonferroni when FDR control would be more appropriate (and vice versa)

Looking Ahead

Next class (Lesson 15): Neyman-Pearson Lemma; Likelihood Ratio Tests (Devore 9.5)

- Most powerful tests for simple hypotheses
- The Neyman-Pearson framework
- Likelihood ratio test statistic
- Connection to tests we've already learned

References

- Devore, Berk, and Carlton. *Modern Mathematical Statistics with Applications* (Springer). Sections 9.2, 9.4
- Chihara and Hesterberg. *Mathematical Statistics with Resampling and R* (Wiley). Sections 8.1–8.2
- Wasserstein, R. L., & Lazar, N. A. (2016). The ASA Statement on p-Values: Context, Process, and Purpose. *The American Statistician*, 70(2), 129–133.
- Fisher, R. A. (1922). On the interpretation of χ^2 from contingency tables, and the calculation of P. *Journal of the Royal Statistical Society*, 85(1), 87–94.
- Benjamini, Y., & Hochberg, Y. (1995). Controlling the false discovery rate: A practical and powerful approach to multiple testing. *Journal of the Royal Statistical Society: Series B*, 57(1), 289–300.

Questions?

Thank you!